

# QuantMinds International

Intercontinental O2 London, November 17, 2025

Rough volatility workshop

Lecture 2: Rough volatility models

Jim Gatheral

Department of Mathematics



**The City University of New York**

## Outline of Lecture 2

- The forward variance curve
- Change of measure
- The rough Bergomi model
- The rough Heston model
- The quadratic rough Heston model

## A simplifying assumption

- We will set interest rates and dividends to zero (unless we specifically say otherwise).
  - It is typically easy to reintroduce nonzero rates and dividends - but of course everything gets more complicated.
- With this assumption
  - The stock price process is a (local) martingale, and so has zero drift.
  - The drift of the stock price process under  $\mathbb{P}$  is the *equity risk premium* - the extra return that investors require for holding the risky stock.

## Stochastic volatility

- Under the pricing measure  $\mathbb{Q}$ , stochastic volatility models take the form

$$\frac{dS_t}{S_t} = \sqrt{V_t} dZ_t$$

with  $V_t = \frac{d}{dt} \langle \log S \rangle_t$ .

- Thus, the stock price process and the quadratic variation process are both assumed continuous.

- There are no jumps!
- To ensure no-arbitrage, the stock price  $S$  is modeled as a positive semimartingale.
  - This excludes for example fractional Brownian motion with  $H \neq \frac{1}{2}$ .
    - If  $H > 1/2$ , quadratic variation (QV) vanishes; if  $H < 1/2$ , QV is infinite!

## Forward variance models

- In classical models, such as Black-Scholes and classical Heston, the volatility process is modeled directly.
  - However  $V_t$  is not observable.
- Bergomi and Guyon<sup>[3]</sup>, suggested that it is natural to model forward variances

$$\xi_t(u) = \mathbb{E}_t^{\mathbb{Q}} [V_u], \quad u > t.$$

- The forward variances, being conditional expectations under  $\mathbb{Q}$ , are tradable.
  - Not only in principle, but in practice, as forward starting variance swaps.
- The forward variance processes are modeled (in the single-factor case) as

$$d\xi_t(u) = \eta_t(u; \omega) dW_t,$$

where as before  $d\langle W, Z \rangle_t = \rho dt$ .

- The  $\mathbb{R}_{\geq 0}$ -valued stochastic process  $\eta_t(u; \omega)$  is progressively measurable for all  $u > 0$ .
  - Conventionally, it is assumed that  $\eta$  is adapted to the filtration generated by  $W$ . In other words, the variance process depends only on the history of the variance process.
  - A truly path-dependent model would have  $\eta$  adapted to the filtration jointly generated by  $W$  and  $Z$ .
- If  $V$  is continuous and uniformly integrable, we can recover  $V_t$  from  $\xi_t(u)$  as  $V_t = \lim_{u \downarrow t} \xi_t(u)$ .
  - For our purposes,  $V_t = \xi_t(t)$ .
- The initial conditions of a typical forward variance model are the initial stock price  $S_t$  and the initial forward variance curve  $\xi_t(u)_{u>t}$ .

## Forward variance curve models and perfect hedging

- As noted by [El Euch and Rosenbaum]<sup>[7]</sup>, models written in forward variance form are explicitly Markovian in the asset price  $S_t$  and the (infinite-dimensional) forward variance curve  $\xi_t$ .
  - European payoffs  $V$  may be perfectly hedged.
  - The delta-hedging strategy involves holding  $\partial_S V$  in the asset and  $\delta_\xi V$  in forward variance contracts where  $\delta_\xi$  denotes the Fréchet derivative of  $V$  with respect to the forward variance curve.

## Example: The Heston model

- The classical Heston model reads:

$$dV_t = -\lambda (V_t - \bar{V}) dt + \nu \sqrt{V_t} dW_t.$$

- The forward variance curve is easily computed as the solution of an ODE:

$$\xi_t(u) := \mathbb{E}_t [V_u] = (V_t - \bar{V}) e^{\lambda(u-t)} + \bar{V}.$$

- Thus, in forward variance form, classical Heston model reads:

$$d\xi_t(u) = \nu e^{\lambda(u-t)} \sqrt{V_t} dW_t.$$

- The classical Heston model generates a term structure of volatility skew  $\mathcal{S}(\tau)$  that is something like

$$\mathcal{S}(\tau) \sim \frac{1}{\lambda\tau} \left\{ 1 - \frac{1 - e^{-\lambda\tau}}{\lambda\tau} \right\}.$$

## Example: The Bergomi model

- In Lecture 1, we demonstrated that smiles typically scale as a power-law.
- Partially motivated by this, Bergomi introduced the  $n$ -factor Bergomi variance curve model:

(1)

$$\xi_t(u) = \xi_0(u) \exp \left\{ \sum_{i=1}^n \eta_i \int_0^t e^{-\kappa_i(u-s)} dW_s^{(i)} + \text{drift} \right\}.$$

- The Bergomi model generates a term structure of volatility skew  $\mathcal{S}(\tau)$  that is something like

$$\mathcal{S}(\tau) \sim \sum_i \frac{1}{\kappa_i \tau} \left\{ 1 - \frac{1 - e^{-\kappa_i \tau}}{\kappa_i \tau} \right\}.$$

- This functional form is related to the term structure of the autocorrelation function.
  - Which is in turn driven by the exponential kernel in the exponent in (1).
- To achieve a decent fit to the observed volatility surface, and to control the forward smile, we need at least two factors.
  - In the two-factor case, there are 8 parameters.
- When calibrating, we find that the two-factor Bergomi model is already over-parameterized. Any combination of parameters that gives a roughly  $1/\sqrt{T}$  ATM skew fits well enough.
  - Moreover, the calibrated correlations between the Brownian increments  $dW_s^{(i)}$  tend to be high.

## ATM skew in the Bergomi model

- The Bergomi model generates a term structure of volatility skew  $\mathcal{S}(\tau)$  that is something like

$$\mathcal{S}(\tau) = \sum_i \frac{1}{\kappa_i \tau} \left\{ 1 - \frac{1 - e^{-\kappa_i \tau}}{\kappa_i \tau} \right\}.$$

- This functional form is related to the term structure of the autocorrelation function.
  - Which is in turn driven by the exponential kernel in the exponent in (1).

## Tinkering with the Bergomi model

- Empirically,  $\mathcal{S}(\tau) \sim \tau^{-\alpha}$  for some  $\alpha$ .

- It's tempting to replace the exponential kernels in (1) with a power-law kernel.

- This would give a model of the form

$$\xi_t(u) = \xi_0(u) \exp \left\{ \eta \int_0^t \frac{dW_s}{(u-s)^\gamma} + \text{drift} \right\}$$

which looks similar to

$$\xi_t(u) = \xi_0(u) \exp \{ \eta W_t^H + \text{drift} \}$$

where  $W_t^H$  is fractional Brownian motion.

## Representations of fBm

There are infinitely many possible representations of fBm in terms of Brownian motion. For example, with  $\gamma = \frac{1}{2} - H$ ,

### Mandelbrot-Van Ness

$$W_t^H = C_H \left\{ \int_{-\infty}^t \frac{dW_s}{(t-s)^\gamma} - \int_{-\infty}^0 \frac{dW_s}{(-s)^\gamma} \right\}.$$

where the choice

$$C_H = \sqrt{\frac{2H\Gamma(3/2-H)}{\Gamma(H+1/2)\Gamma(2-2H)}}$$

ensures that

$$\mathbb{E}[W_t^H W_s^H] = \frac{1}{2} \left\{ t^{2H} + s^{2H} - |t-s|^{2H} \right\}.$$

## The RFSV model

In Lecture 1, our analysis of realized variance data suggested the following model for volatility under the real (or historical or physical) measure  $\mathbb{P}$ :

$$\log \sigma_u - \log \sigma_t = \nu (W_u^H - W_t^H), \quad u > t.$$

Then, with the Mandelbrot-Van Ness representation of fBm,

$$\begin{aligned} \log V_u - \log V_t &= 2\nu C_H \left\{ \int_t^u \frac{1}{(u-s)^\gamma} dW_s^{\mathbb{P}} \right. \\ &\quad \left. + \int_{-\infty}^t \left[ \frac{1}{(u-s)^\gamma} - \frac{1}{(t-s)^\gamma} \right] dW_s^{\mathbb{P}} \right\} \\ &=: 2\nu C_H [M_t(u) + Z_t(u)]. \end{aligned}$$

- Note that  $\mathbb{E}_t^{\mathbb{P}} [M_t(u)] = 0$  and  $Z_t(u)$  is  $\mathcal{F}_t$ -measurable.
  - To price options, it would seem that we would need to know  $\mathcal{F}_t$ , the entire history of the Brownian motion  $W_s$  for  $s < t$ !

## Pricing under $\mathbb{P}$

- Let  $\tilde{\eta} = 2\nu C_H$  with  $\tilde{\eta} = \eta\sqrt{2H}$ . Then

$$\begin{aligned} V_u &= V_t \exp \left\{ \tilde{\eta} \int_t^u \frac{1}{(u-s)^\gamma} dW_s^{\mathbb{P}} + 2\nu C_H Z_t(u) \right\} \\ &= \mathbb{E}_t^{\mathbb{P}} [V_u] \mathcal{E} \left( \tilde{\eta} \int_t^u \frac{1}{(u-s)^\gamma} dW_s^{\mathbb{P}} \right). \end{aligned}$$

- The conditional distribution of  $V_u$  depends on  $\mathcal{F}_t$  only through the variance forecasts  $\mathbb{E}_t^{\mathbb{P}} [V_u]$ ,
- The last equality is the key:
  - To price options, one does not need to know  $\mathcal{F}_t$ , the entire history of the Brownian motion  $W_s^{\mathbb{P}}$  for  $s < t$ .

## Pricing under $\mathbb{Q}$

- Our model under  $\mathbb{P}$  reads:

$$V_u = \mathbb{E}_t^{\mathbb{P}} [V_u] \mathcal{E} \left( \tilde{\eta} \int_t^u \frac{1}{(u-s)^\gamma} dW_s^{\mathbb{P}} \right).$$

- Consider some general change of measure

$$dW_s^{\mathbb{P}} = dW_s^{\mathbb{Q}} + \lambda_s ds,$$

where  $\{\lambda_s : s > t\}$  has a natural interpretation as the price of volatility risk. We may then write

$$V_u = \mathbb{E}_t^{\mathbb{P}} [V_u] \mathcal{E} \left( \tilde{\eta} \int_t^u \frac{1}{(u-s)^\gamma} dW_s^{\mathbb{P}} \right) \exp \left\{ \tilde{\eta} \int_t^u \frac{\lambda_s}{(u-s)^\gamma} ds \right\}.$$

- Although the conditional distribution of  $V_u$  under  $\mathbb{P}$  is lognormal, it will not be lognormal in general under  $\mathbb{Q}$ .
  - The upward sloping smile in VIX options means  $\lambda_s$  cannot be deterministic in this picture.

## The rough Bergomi model

Let's nevertheless consider the simplest change of measure

$$dW_s^{\mathbb{P}} = dW_s^{\mathbb{Q}} + \lambda(s) ds,$$

where  $\lambda(s)$  is a deterministic function of  $s$ . Then from (2), we would have

$$\begin{aligned} V_u &= \mathbb{E}_t^{\mathbb{P}} [V_u] \mathcal{E} \left( \eta \tilde{W}_t^{\mathbb{Q}}(u) \right) \exp \left\{ \tilde{\eta} \int_t^u \frac{1}{(u-s)^\gamma} \lambda(s) ds \right\} \\ &= \xi_t(u) \mathcal{E} \left( \eta \tilde{W}_t^{\mathbb{Q}}(u) \right) \end{aligned} \tag{1}$$

where the forward variances  $\xi_t(u) = \mathbb{E}_t^{\mathbb{Q}} [V_u]$  are (at least in principle) tradable and observed in the market.

- $\xi_t(u)$  is the product of two terms:
- $\mathbb{E}_t^{\mathbb{P}} [V_u]$  which depends on the historical path  $\{W_s, s < t\}$  of the Brownian motion

- a term which depends on the price of risk  $\lambda(s)$ .

## Features of the rough Bergomi model

- The rBergomi model is a non-Markovian generalization of the Bergomi model:

$$\mathbb{E}[V_u | \mathcal{F}_t] \neq \mathbb{E}[V_u | V_t].$$

- The rBergomi model is Markovian in the (infinite-dimensional) state vector  $\mathbb{E}^{\mathbb{Q}}[V_u | \mathcal{F}_t] = \xi_t(u)$ .

- We have achieved our earlier aim of replacing the exponential kernels in the Bergomi model with a power-law kernel.
- We may therefore expect that the rBergomi model will generate a realistic term structure of ATM volatility skew.

## Re-interpretation of the conventional Bergomi model

- A conventional  $n$ -factor Bergomi model is not self-consistent for an arbitrary choice of the initial forward variance curve  $\xi_t(u)$ .
- $\xi_t(u) = \mathbb{E}_t[V_u]$  should be consistent with the assumed dynamics.
- Viewed from the perspective of the fractional Bergomi model however:
  - The initial curve  $\xi_t(u)$  reflects the history  $\{W_s; s < t\}$  of the driving Brownian motion up to time  $t$ .
  - The exponential kernels in the exponent of the conventional Bergomi model approximate more realistic power-law kernels.
- The conventional two-factor Bergomi model is then justified in practice as a tractable Markovian engineering approximation to a more realistic fractional Bergomi model.

## The variance contract and the log-contract.

Under zero interest rates and dividends, applying Itô's Lemma, path-by-path

$$\log\left(\frac{S_T}{S_t}\right) = \int_t^T d\log(S_u) = \int_t^T \frac{dS_u}{S_u} - \frac{1}{2} \int_t^T V_u du.$$

- The second term on the RHS is immediately recognizable as half the quadratic variation  $\langle X \rangle_T$  over the interval  $[0, T]$ .
- Thus, the value of the variance contract

$$M_t(T) := \mathbb{E}_t \left[ \int_t^T V_u du \right] = -2 \mathbb{E}_t \left[ \log \frac{S_T}{S_t} \right],$$

may be expressed in terms of the fair value of the log-contract.

## Robust valuation of the variance contract

- The log-contract may be valued using the Carr-Madan spanning formula as the *log-strip* of options that gives rise to the VIX formula.

- In principle, we need to know the prices of Europeans with all possible strikes for a given expiration  $T$ .
  - In practice, we only have a finite number of strike prices listed per expiration.
- One way to estimate the value of such swaps would be to fit a parameterization such as SVI or one of the [Vola Dyanamics](#) curves, interpolating and extrapolating to fill in all the other strikes.
- We will now show how to estimate the value of the variance contract robustly with not too much dependence on the interpolation/extrapolation method.

## A cool formula

- Write  $\Sigma(k, T) = \Sigma$  for short and define

$$d_{\pm} = -\frac{k}{\sqrt{\Sigma}} \pm \frac{\sqrt{\Sigma}}{2}$$

- Further define the inverse functions  $g_{\pm}(z) = d_{\pm}^{-1}(z)$ .
  - Intuitively,  $z$  measures the log-moneyness of an option in implied standard deviations.
- Then,

$$M_t(T) = -2 \mathbb{E}_t \left[ \log \frac{S_T}{S_t} \right] = \int_{-\infty}^{\infty} dz N'(z) \Sigma(g_-(z)).$$

- To see this formula is plausible, it is obviously correct in the flat-volatility Black-Scholes case.

## Estimating the forward variance curve in practice

- With the above formulae, it's easy to see how to get the forward variance curve in principle.
- Let's now do this in practice.
- Again, we consider the SPX volatility surface as of 15-Feb-2023.

## The SPX volatility surface as of 15-Feb-2023

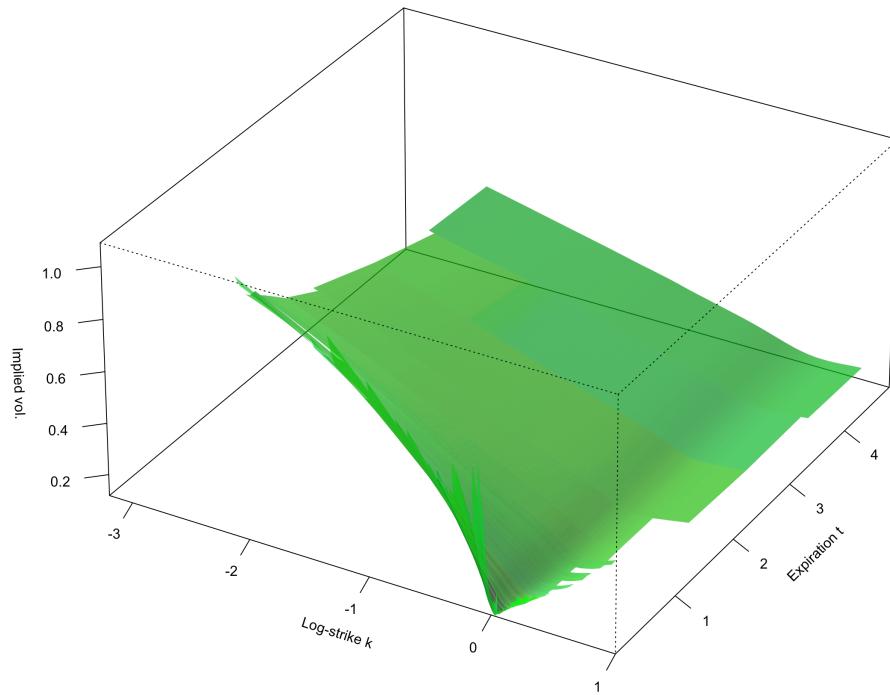


Figure 1. The SPX volatility surface as of 15-Feb-2023 (data from OptionMetrics via WRDS).

## Set up the environment

```
In [1]: # Download .zip from GitHub
release_url <- "https://github.com/jgatheral/RoughVolatilityWorkshop2025/raw/main/QM_2025.zip"
local_zip <- "QM_2025.zip"

download.file(release_url, local_zip, mode = "wb")
unzip(local_zip)
```

```
In [2]: library(repr)
library(colorspace)
library(stinepack)
library(foreach)
library(doParallel)
```

Loading required package: iterators

Loading required package: parallel

## Some R-code

```
In [3]: source("black_scholes.R")
source("black_formula.R")
source("fwd_var_curve.R")
source("fukasawa_robust.R")
source("gamma_kernel.R")
source("hybrid_bss.R")
source("lewis.R")
source("plot_ivols.R")
```

```
source("plot_ivols_mc.R")
source("rough_heston_pade.R")
```

## Load volatility smiles from 15-Feb-2023

```
In [4]: load("spx_vix_vols_20230215.rData")
ivolData <- spxIvols20230215
head(ivolData)
```

A data.frame: 6 × 7

	Expiry	Texp	Strike	Bid	Ask	Fwd	CallMid
				<int>	<dbl>	<dbl>	<dbl>
1	20230216	0.002737851	1000	NA	7.793085	4146.742	NA
2	20230216	0.002737851	1200	NA	6.813266	4146.742	NA
3	20230216	0.002737851	1400	NA	5.987566	4146.742	NA
4	20230216	0.002737851	1600	NA	5.273554	4146.742	NA
5	20230216	0.002737851	1800	NA	4.644049	4146.742	NA
6	20230216	0.002737851	2000	NA	4.080578	4146.742	NA

```
In [5]: options(repr.plot.width=10, repr.plot.height=7, repr.plot.res=150)
res.plot <- plotIvols(ivolData)
```

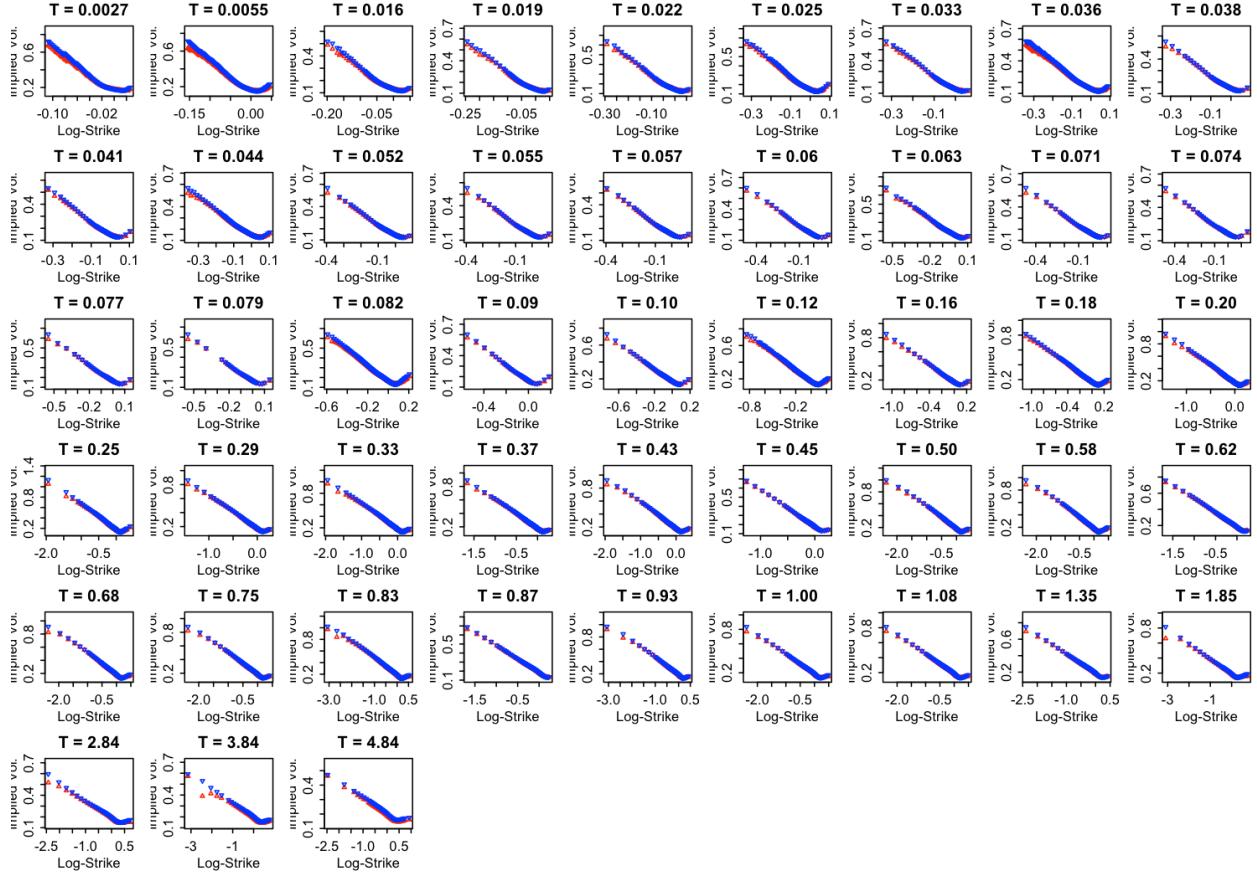


Figure 2: SPX smiles as of February 15, 2023.

## Set up nice colors

```
In [6]: my.col <- sequential_hcl(5, palette="Batlow")
bl <- "royalblue"
rd <- "red2"
pk <- "hotpink1"
gr <- "green4"
br <- "brown"
pu <- "purple"
or <- "orange"
```

## Robust estimation of the variance contract

```
In [7]: expiries <- res.plot$expiries
(vs <- varSwap.Robust(ivolData)$vs.mid)
```

```
0.036529328507355 · 0.0317776298748159 · 0.019801436839558 · 0.0216205797598485 ·
0.0239817142815479 · 0.0260070933624724 · 0.0230480008871306 · 0.0242269111404731 ·
0.0254217621840437 · 0.0262608820432924 · 0.0272812740956352 · 0.0251873317922458 ·
0.0270816758954462 · 0.0277890020550951 · 0.0283828078588884 · 0.0307845544704758 ·
0.0288281365209749 · 0.0326234863276105 · 0.033036923041284 · 0.0333085545046597 ·
0.0333045778294113 · 0.0321547119979983 · 0.0372856381939817 · 0.0368526562667083 ·
0.0384015885110832 · 0.0382728883769698 · 0.0389762987224796 · 0.0422129823407889 ·
0.041766324573276 · 0.0449034147120543 · 0.0453292487617555 · 0.0467388646631131 ·
0.046384306264404 · 0.0493899123530733 · 0.0506716850375437 · 0.0509751148828293 ·
0.0525838020688622 · 0.0536616949197811 · 0.0552555914782415 · 0.0528434116744946 ·
0.0544748623553668 · 0.0544544038189155 · 0.0549311335723502 · 0.0552749722502056 ·
0.0571945373055973 · 0.0567880111485679 · 0.0571112166824289 · 0.059465108018572
```

```
In [8]: plot(expiries, vs, type="b", pch=20, lwd=2, col=pk, ylab="Swap values", xlab="Maturity (years)", log='
```

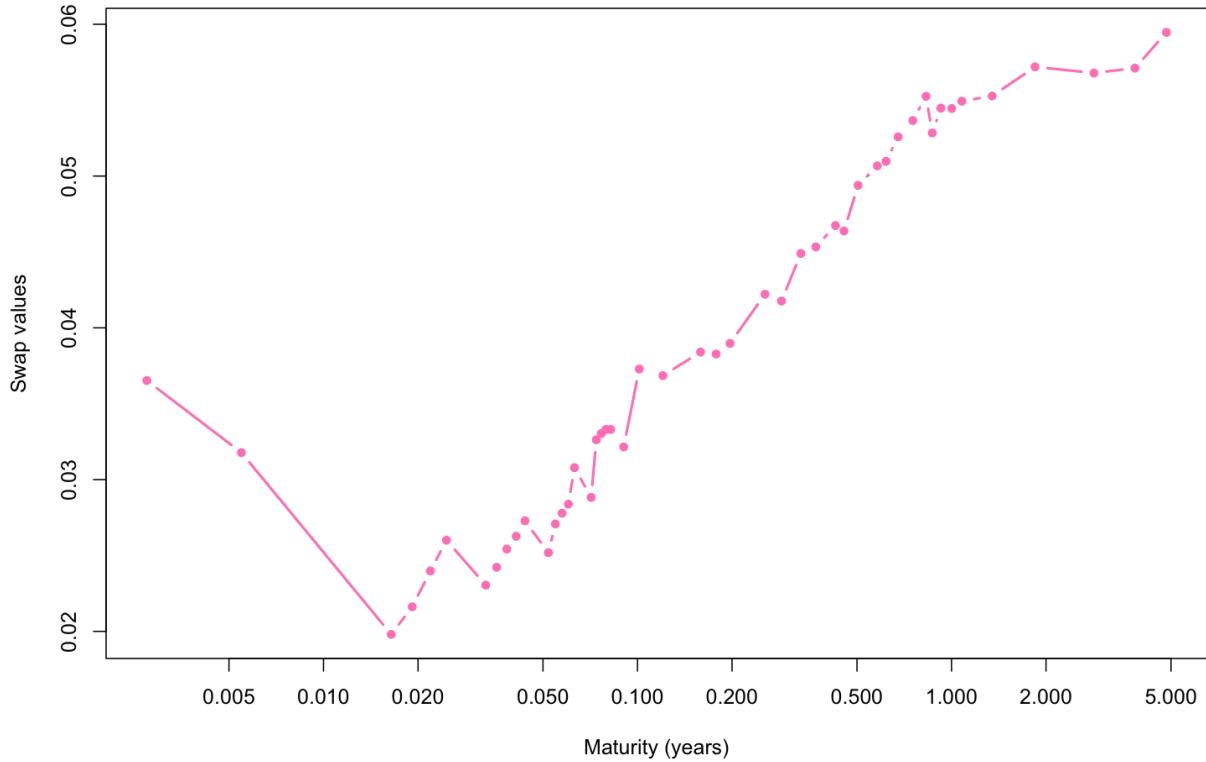


Figure 3: Log-linear plot of the forward variance curve.

## The forward variance curve from the variance curve

- By definition,  $\xi_t(u) = \mathbb{E}_t[V_u]$ .
- Recall that

$$M_t(T) = \int_t^T \xi_t(u) du.$$

- Differentiating wrt  $T$  gives

$$\xi_t(T) = \partial_T M_t(T).$$

## Exact smooth forward curve construction

- In 2019, Baruch MFE student Rick Cao implemented (beautifully) a beautiful paper of [Filipović and Willems]<sup>[5]</sup>
- That paper presents a non-parametric method to estimate the discount curve from market quotes, that reproduces the market quotes perfectly and has maximal smoothness in the sense that it minimizes the  $L^2$ -norm of the forward curve.
- We apply this method to the variance swap curve.
  - The resulting forward variance curve is piecewise quadratic.
- Warning: The resulting forward variance curve is not guaranteed to be positive - though this does not seem to matter in practice.

## Adding a bid-offer spread to smooth the curve

- In practice, `w.in` is not known exactly but only up to some bid-offer spread.
  - Moreover some expirations have more strikes than others and we get a better estimate.
- We input this bid-offer volatility spread using the `eps` parameter.

## Why the forward variance curve should be smooth

- If there are two forward variance curves that are consistent with the data, the smoother one is better.
- To see why, consider trading forward variance swaps around discontinuities in the forward curve!

```
In [9]: xi.curve.smooth
```

```

function (expiries, w.in, xi = TRUE, eps = 0)
{
  phi <- function(tau) {
    function(x) {
      min <- pmin(x, tau)
      return(1 - min^3/6 + x * tau * (2 + min)/2)
    }
  }
  phi.deri <- function(tau) {
    function(x) {
      min <- pmin(x, tau)
      return(tau - min^2/2 + tau * min)
    }
  }
  n <- length(expiries)
  c <- diag(n)
  A <- array(dim = c(n, n))
  for (i in seq(1, n)) {
    A[i, ] <- phi(expiries[i])(expiries)
  }
  A.inv <- solve(A)
  obj.1 <- function(err.vec) {
    v <- w.in + 2 * sqrt(w.in) * err.vec * sqrt(expiries)
    return(t(v) %*% A.inv %*% v)
  }
  res.optim <- optim(rep(0, n), obj.1, method = "L-BFGS-B",
    lower = rep(-eps, n), upper = rep(eps, n))
  err.vec <- res.optim$par
  w.in.1 <- w.in + 2 * sqrt(w.in) * err.vec * sqrt(expiries)
  Z <- A.inv %*% w.in.1
  curve.raw <- function(x) {
    sum.curve <- 0
    sum.curve.deri <- 0
    for (i in seq(1, n)) {
      sum.curve <- sum.curve + Z[i] * phi(expiries[i])(x)
      sum.curve.deri <- sum.curve.deri + Z[i] * phi.deri(expiries[i])(x)
    }
    if (xi) {
      return(sum.curve.deri)
    }
    else {
      return(sum.curve)
    }
  }
  xi.curve.out <- Vectorize(curve.raw)
  fit errs <- sqrt(w.in.1/expiries) - sqrt(w.in/expiries)
  return(list(xi.curve = xi.curve.out, fit errs = fit errs,
    w.out = w.in.1))
}

```

```
In [10]: w.in <- vs*expiries
xi.smooth <- xi.curve.smooth(expiries, w.in, eps = .006)
xi.smooth.c <- xi.smooth$xi.curve
xi.smooth.w.out <- xi.smooth$w.out

In [11]: curve(xi.smooth.c, from=0, to=3, col=bl, lty=1, lwd=2, xlab="Maturity", ylab="Forward variance", n=100)
```

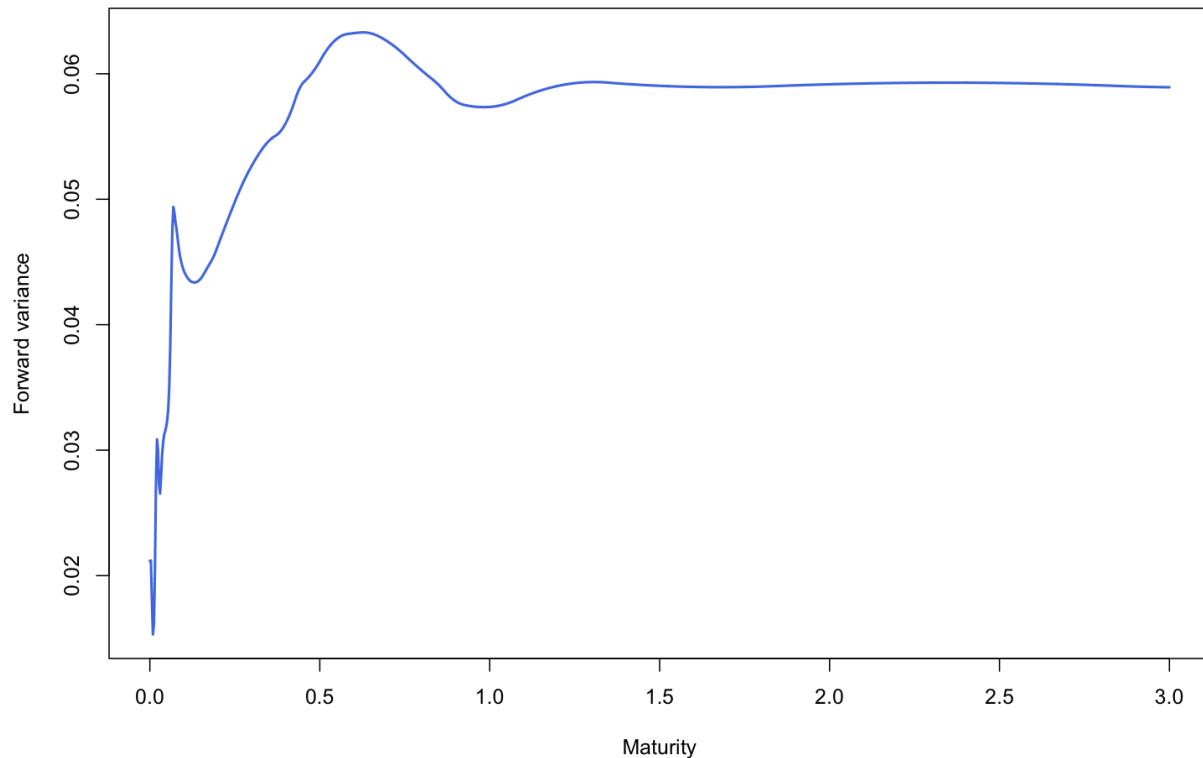


Figure 4: Smooth approximation to the forward variance curve.

### What does exact interpolation mean?

- Output variance swaps exactly match input variance swaps (up to the bid-offer spread).

```
In [12]: plot(expiries, w.in/expiries, col=pk, lwd=2, pch=20, xlab="Maturity", ylab="Variance rate", type="b")
points(expiries, xi.smooth.w.out/expiries, col=bl, pch=20, lwd=2, type="b")
```

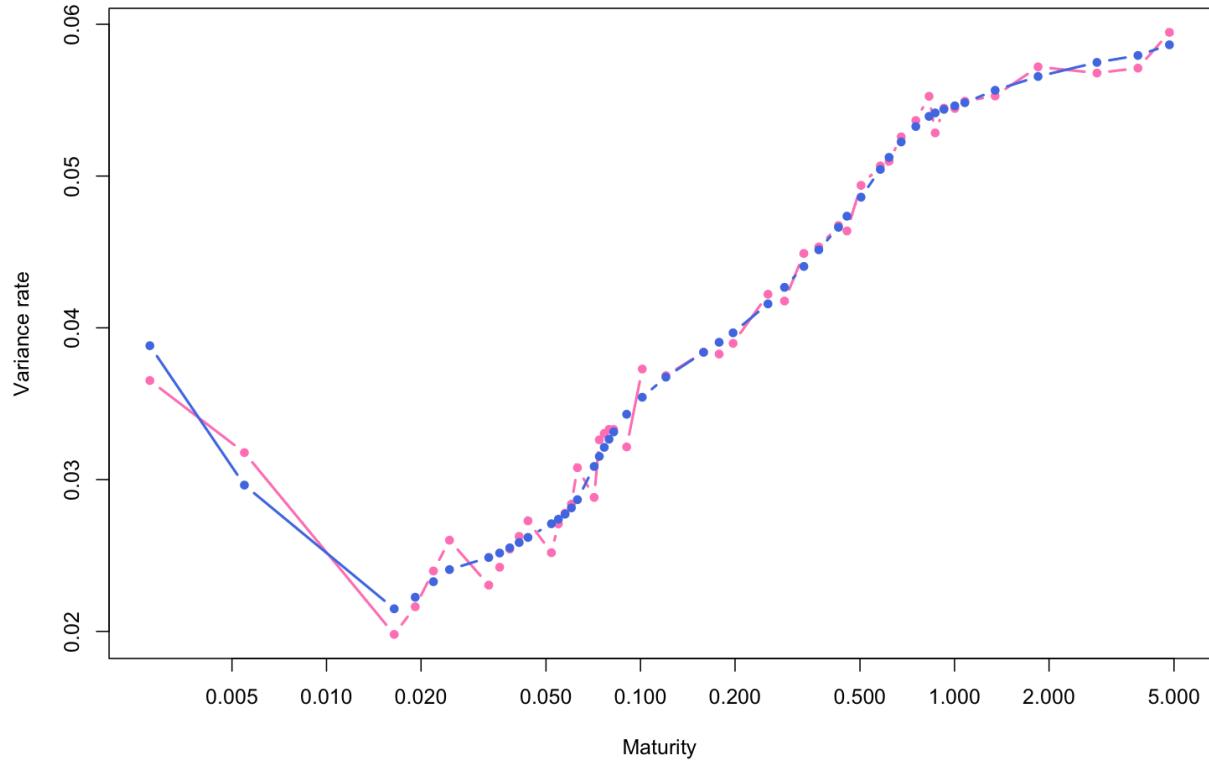


Figure 5: Blue circles are input variance swaps; pink circles are output variance swaps; green dots are from variance swaps with bid/offer spread.

## Instantiate the forward variance curve

We can speed things up by instantiating the curve:

```
In [13]: xiCurveObj <- CurveSmoothBuilder(expiries, vs*expiries, eps=.006)
xiCurveObj$fitCurve()
xi.curve.fast <- xiCurveObj$getForwardVarCurve()
```

```
In [14]: save(xiCurveObj, file="xi20230215.rData")
```

Note: Object-oriented programming in R. We can even save the object for future use!

## Plot the fast instantiated forward variance curve

```
In [15]: curve(xi.smooth.c, from=0, to=3, col=bl, lty=1, lwd=2, xlab="Maturity", ylab="Forward variance")
curve(xi.curve.fast, from=0, to=3, col=pk, lwd=3, add=T, lty=2)
```

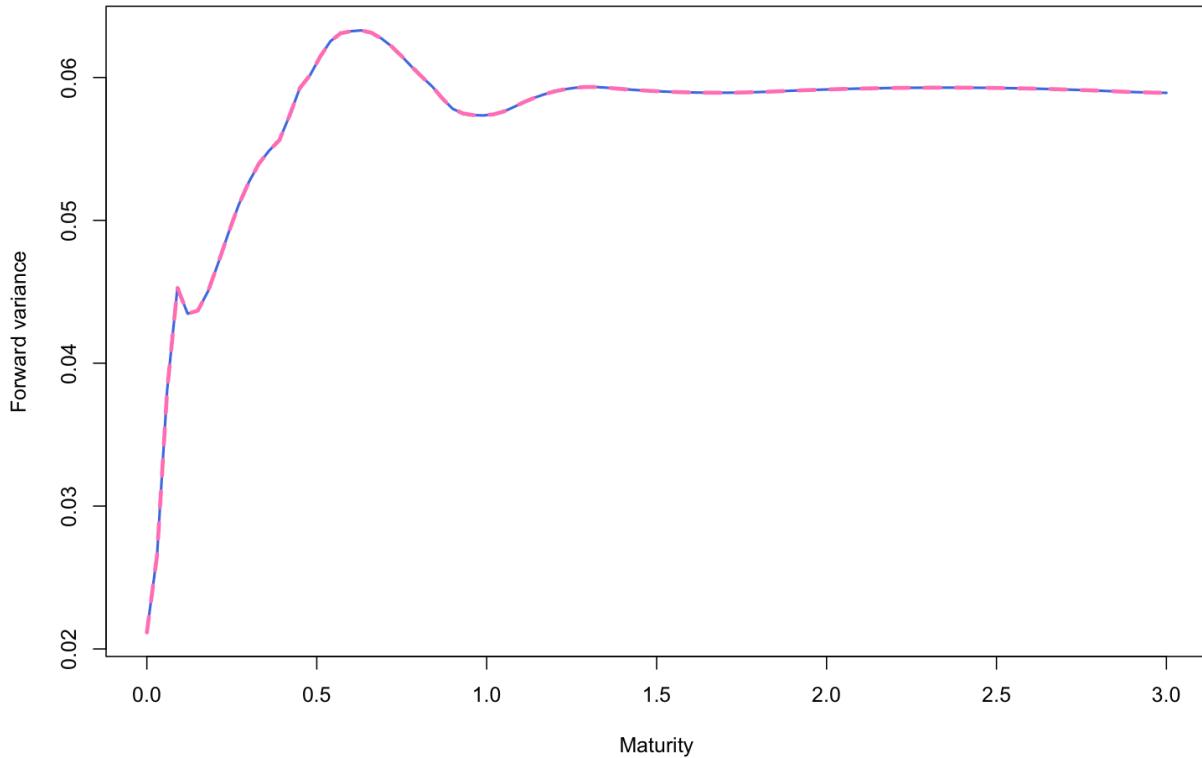


Figure 6: Original and fast (instantiated) forward variance curve.

## Using forward variance curves in practice

- Forward variance models, rough volatility models in particular, take the forward variance curve as given.
- The forward variance curve is estimated from the variance swap curve (by differencing for example).
  - Variance swap estimates depend on the extrapolation methodology.
- A forward variance model that takes the estimated forward variance curve as input will not generate the same variance swap values as the estimates.
  - In particular, the smile extrapolation will be different.
- In practice therefore, we iterate on the forward variance curve so as to match market and model ATM volatilities.

## The stock price process

- The observed anticorrelation between price moves and volatility moves may be modeled naturally by anticorrelating the Brownian motion  $W$  that drives the volatility process with the Brownian motion driving the price process.
- Thus

$$\frac{dS_t}{S_t} = \sqrt{V_t} dZ_t$$

with

$$dZ_t = \rho dW_t + \sqrt{1 - \rho^2} dW_t^\perp$$

where  $\rho$  is the correlation between volatility moves and price moves.

## Hybrid simulation of BSS processes

- The Rough Bergomi variance process is a special case of a Brownian Semistationary (BSS) process.
- [Bennedsen, Lunde and Pakkanen]<sup>[3]</sup> show how to simulate such processes more efficiently.
- [McCrickerd and Pakkanen]<sup>[10]</sup> show how to increase the efficiency of the hybrid scheme with variance reduction.
  - Moreover, they provide a sample Jupyter notebook!
- An improved version of their idea is roughly as follows:

$$\begin{aligned} \int_t^u \frac{dW_s}{(u-s)^\gamma} &= \int_0^\tau \frac{dW_s}{s^\gamma} \\ &\approx \sum_{k=1}^{\kappa} \int_{t_{k-1}}^{t_k} \frac{dW_s}{s^\gamma} + \sum_{k=\kappa+1}^n b_k \int_{t_{k-1}}^{t_k} dW_s \\ &= \sum_{k=1}^{\kappa} \int_{t_{k-1}}^{t_k} \frac{dW_s}{s^\gamma} + \sum_{k=\kappa+1}^n b_k Z_k \sqrt{\Delta} \end{aligned}$$

where  $\gamma = \frac{1}{2} - H$ ,  $\Delta = (u-t)/n$ ,  $t_k = k\Delta$ ,

- $b_k^2 = \text{var} \left[ \int_{t_{k-1}}^{t_k} \frac{dW_s}{s^\gamma} \right] = \int_{t_{k-1}}^{t_k} \frac{ds}{s^{2\gamma}} = \frac{1}{2H} \{ t_k^{2H} - t_{k-1}^{2H} \}$

and the  $Z_k$  are iid  $N(0, 1)$  random variables.

- The choice  $\kappa = 0$  corresponds to the Euler scheme (or Riemann sum scheme), which performs relatively poorly.
- The choice  $\kappa = 1$  works well in practice.
  - The idea is to not only match the variance at each step (which the Riemann sum scheme does) but also the covariance

$$\text{cov} \left[ \int_0^\Delta \frac{dW_s}{s^\gamma}, \int_0^\Delta dW_s \right] = \int_0^\Delta \frac{ds}{s^\gamma} = \frac{\Delta^{1-\gamma}}{1-\gamma} = \frac{\Delta^{H+\frac{1}{2}}}{H+\frac{1}{2}}.$$

- We simulate another normal random variable to achieve this variance and covariance.

## R-implementation of the hybrid scheme

The following function simulates the Riemann-Liouville process

$$\tilde{W}_t^H := \sqrt{2H} \int_0^t (t-s)^{\alpha-1} dW_s,$$

where  $\alpha = H + \frac{1}{2}$ .

In [16]: WtildeRL.sim

```

function (params, hybrid = T)
function(W, Wperp) {
  library(stats)
  steps <- dim(W)[1]
  N <- dim(W)[2]
  stopifnot(dim(Wperp) == c(steps, N))
  dt <- 1/steps
  wp <- Vectorize(wRL(params))
  sqrt.dt <- sqrt(dt)
  tj <- (1:steps) * dt
  wpj <- c(0, wp(tj))
  bstar <- sqrt(diff(wpj)/dt)
  cstar <- cov1RL(params)(dt)
  rhostar <- cstar/(bstar[1] * bstar[2])
  rhobarstar <- sqrt(1 - rhostar^2)
  f <- function(n) {
    Wr <- W[steps:1, n]
    Y.Euler <- convolve(bstar, Wr, type = "open")[1:steps]
    Y.Correct <- bstar[1] * ((rhostar - 1) * W[, n] + rhobarstar *
      Wperp[, n])
    return((Y.Correct * isTRUE(hybrid) + Y.Euler) * sqrt(dt))
  }
  Wtilde <- sapply(1:N, f)
  return(Wtilde)
}

```

The following code uses the same Riemann-Liouville process for each expiry.

In [17]: hybridSchemeRL.S

```

function (params, xi)
function(paths, steps, expiries) {
  eta <- params$eta
  H <- params$al - 1/2
  rho <- params$rho
  N <- paths
  W <- matrix(rnorm(N * steps), nrow = steps, ncol = N)
  Wperp <- matrix(rnorm(N * steps), nrow = steps, ncol = N)
  Zperp <- matrix(rnorm(N * steps), nrow = steps, ncol = N)
  Z <- rho * W + sqrt(1 - rho * rho) * Zperp
  Wtilde <- WtildeRL.sim(params)(W, Wperp)
  sim <- function(expiry) {
    dt <- expiry/steps
    ti <- (1:steps) * dt
    xi.t <- xi(ti)
    v1 <- xi.t * exp(eta * expiry^H * Wtilde - 1/2 * eta^2 *
      ti^(2 * H))
    v0 <- rep(xi(0), N)
    v <- rbind(v0, v1[-steps, ])
    logs <- apply(sqrt(v * dt) * Z - v/2 * dt, 2, sum)
    s <- exp(logs)
    return(s)
  }
  st <- t(sapply(expiries, sim))
  return(st)
}

```

## Run the hybrid BSS scheme

We will use R parallel processing functionality.

```
In [18]: paths <- 1e5
steps <- 200
```

Parameters are from the fit to SPX smiles as of 14-Aug-2013 reported in

```
In [19]: params.rBergomi <- list(al=0.55, eta=2.3, rho=-0.9)
xiCurve <- xi.curve.fast # We use the instantiated curve for speed
```

```
In [20]: t0<-proc.time()

#number of iterations
iters<- max(1,floor(paths/1000))

#setup parallel backend
cl.num <- detectCores() # This number is 8 on my MacBook Pro
cl<-makeCluster(cl.num)
registerDoParallel(cl)

#loop
ls <- foreach(icount(iters),.packages = "stinepack") %dopar% {
  hybridSchemeRL.S(params.rBergomi,xiCurve)(paths=1000, steps=steps, expiries=expiries)
}

stopCluster(cl)
mcMatrix2013 <- do.call(cbind, ls) #Bind all of the submatrices into one big matrix
```

```
print(proc.time() - t0)
```

user system elapsed  
0.173 0.075 8.341

## Plot actual and rough Bergomi (2013) smiles

```
In [21]: system.time(res.plot2013 <- plotIvolsMC(ivolData, mcMatrix=mcMatrix2013))
```

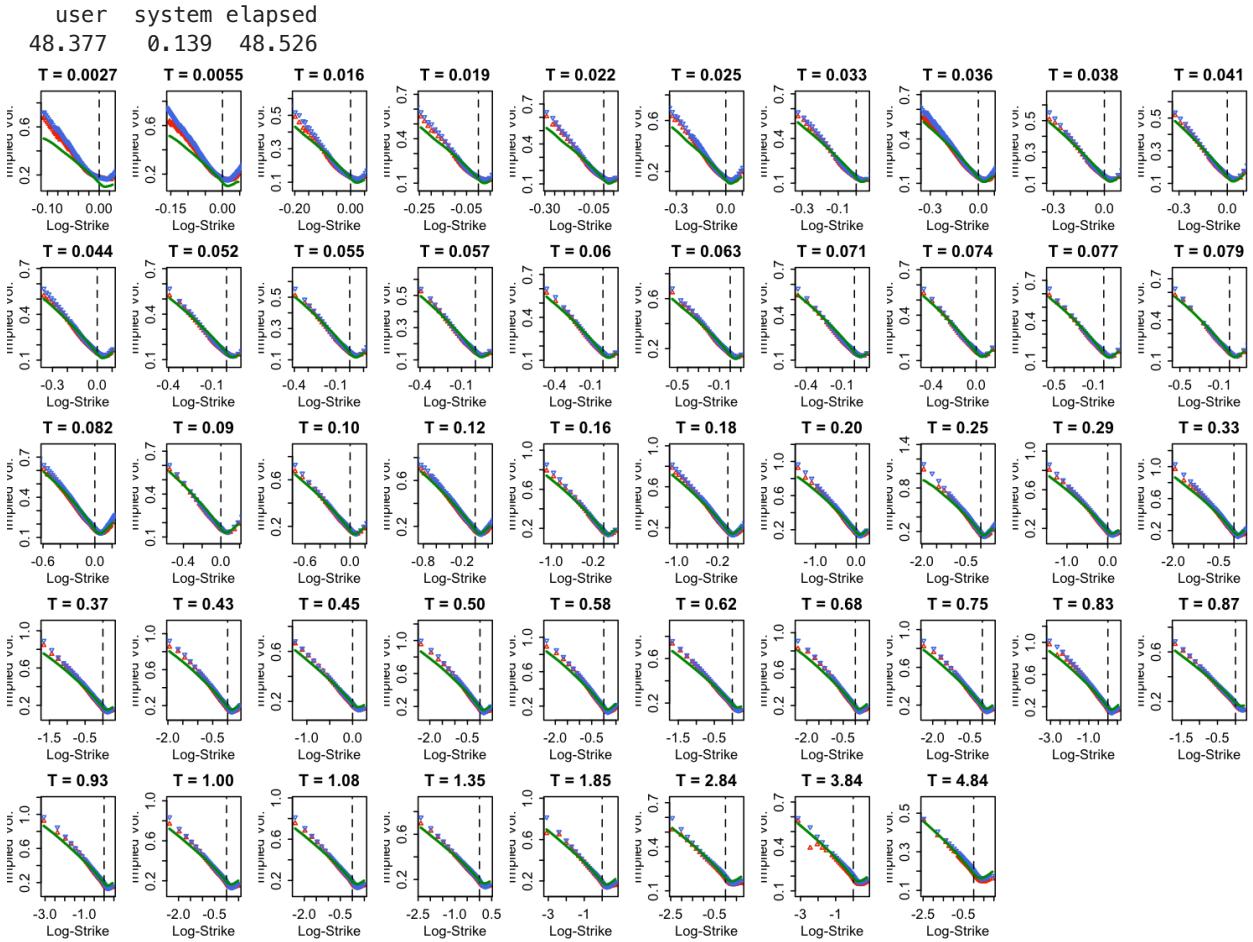


Figure 7: Rough Bergomi smiles (green) with parameters from 2013 superimposed on February 15, 2023 SPX smiles.

## Plot a selection of actual and rough Bergomi smiles

```
In [22]: res.plot2013.6 <- plotIvolsMC(ivolData, mcMatrix=mcMatrix2013, slices=c(2, 10, 21, 28, 34, 42))
```

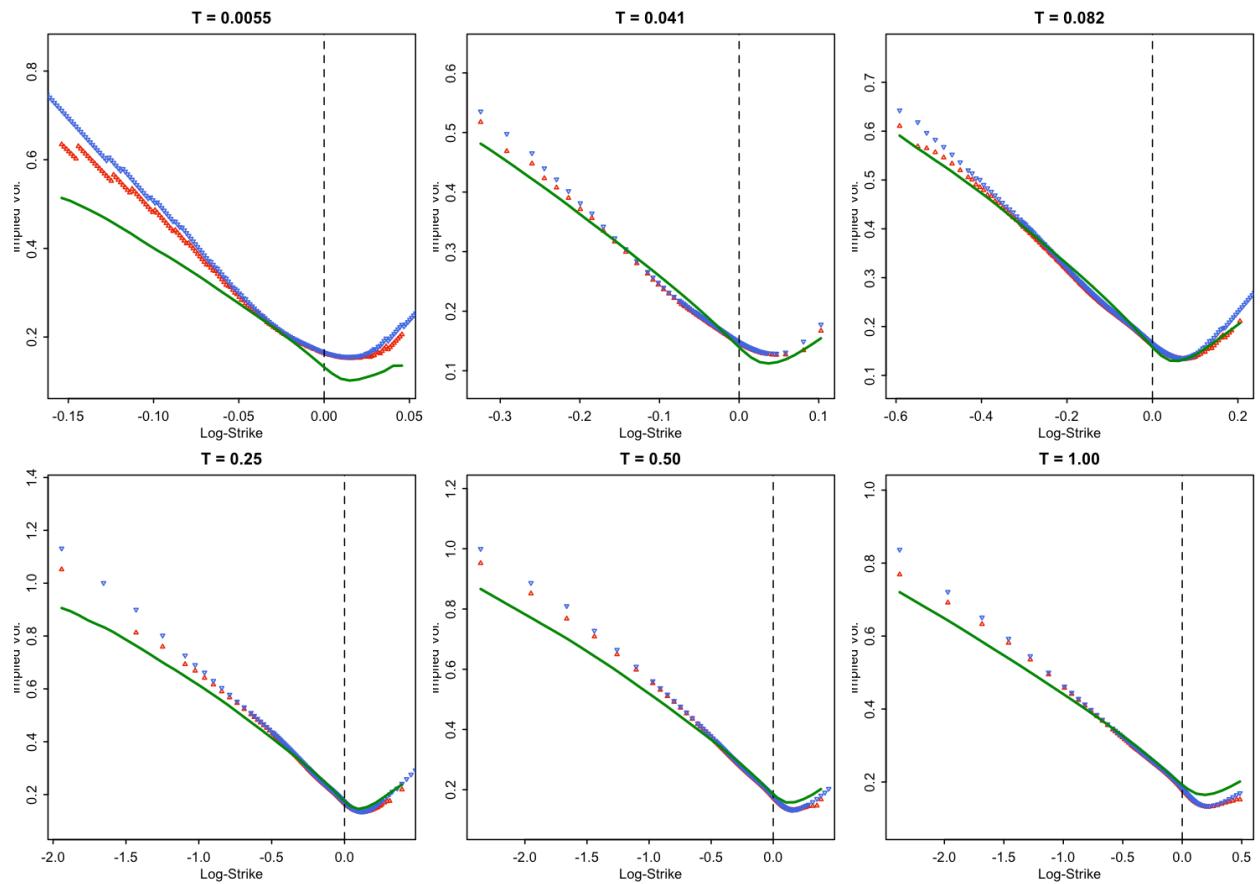


Figure 8: Six rough Bergomi smiles (green) with parameters from 2013 superimposed on February 15, 2023 SPX smiles.

## Comments on Figures 7 and 8

- Considering that we are using parameters from 2013, the rough Bergomi smiles look pretty good.
  - Rough Bergomi parameters seem to be remarkably stable!
- If simulation were fast enough, we could just iterate on these parameters to find the best fit to observed option prices.
  - The BSS scheme is not yet fast enough, at least in my R implementation.

## Plot the ATM skew vs ATM rough Bergomi skew

```
In [23]: plot(res.plot2013$expiries, res.plot2013$atmSkew, pch=20, col=bl, xlab="Expiry", ylab="ATM skew")
lines(res.plot2013$expiries, res.plot2013$atmSkewMC, col=gr, lwd=2)
```

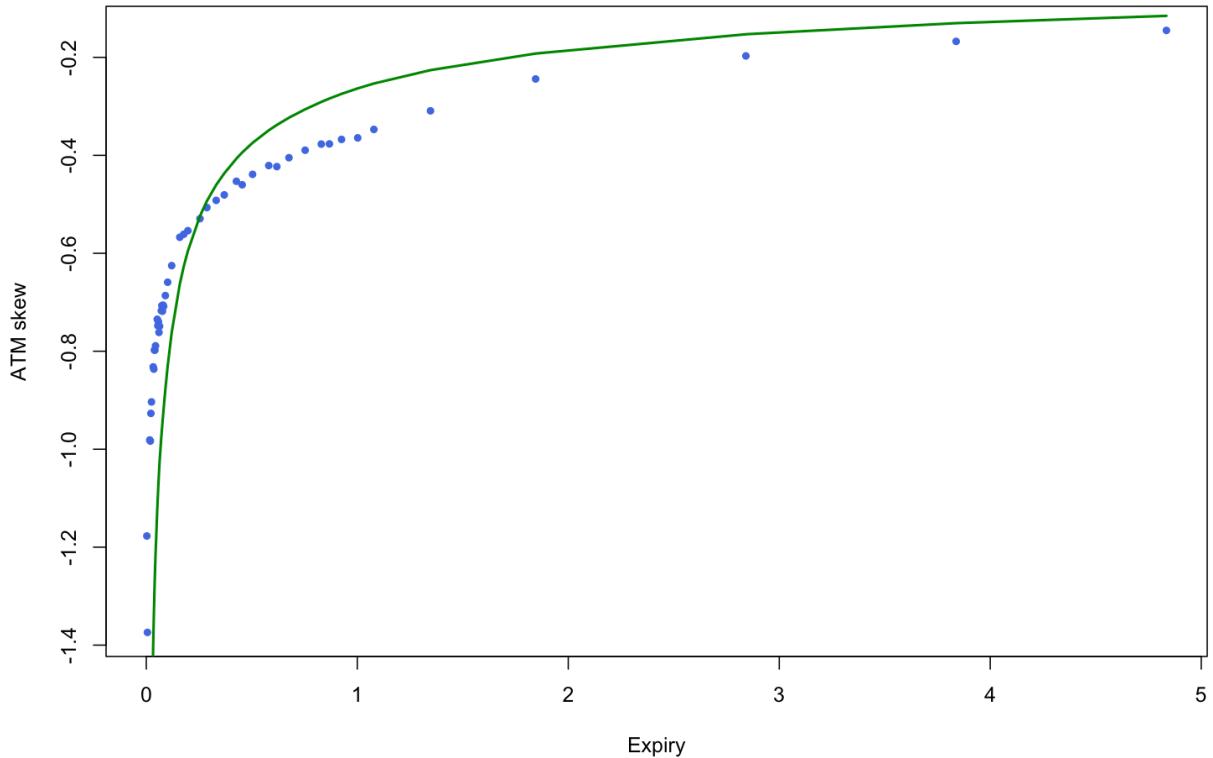


Figure 9: Termstructure of skew seems to be off. Maybe we can do better?

### Guessing rBergomi model parameters

- The rBergomi model has only three parameters:  $H$ ,  $\eta$  and  $\rho$ .
- The model parameters  $H$ ,  $\eta$  and  $\rho$  have very direct interpretations:
  - $H$  controls the decay of ATM skew  $S(\tau)$  for very short expirations.
  - The product  $\rho\eta$  sets the level of the ATM skew for longer expirations.
    - Keeping  $\rho\eta$  constant but decreasing  $\rho$  (so as to make it more negative) pushes the minimum of each smile towards higher strikes.
- So we can guess parameters in practice.
  - A couple of examples of the results of guessing are given in [Bayer, Friz and Gatheral]<sup>[1]</sup>.

### Log-log plot of rough Bergomi ATM skew for various $H$

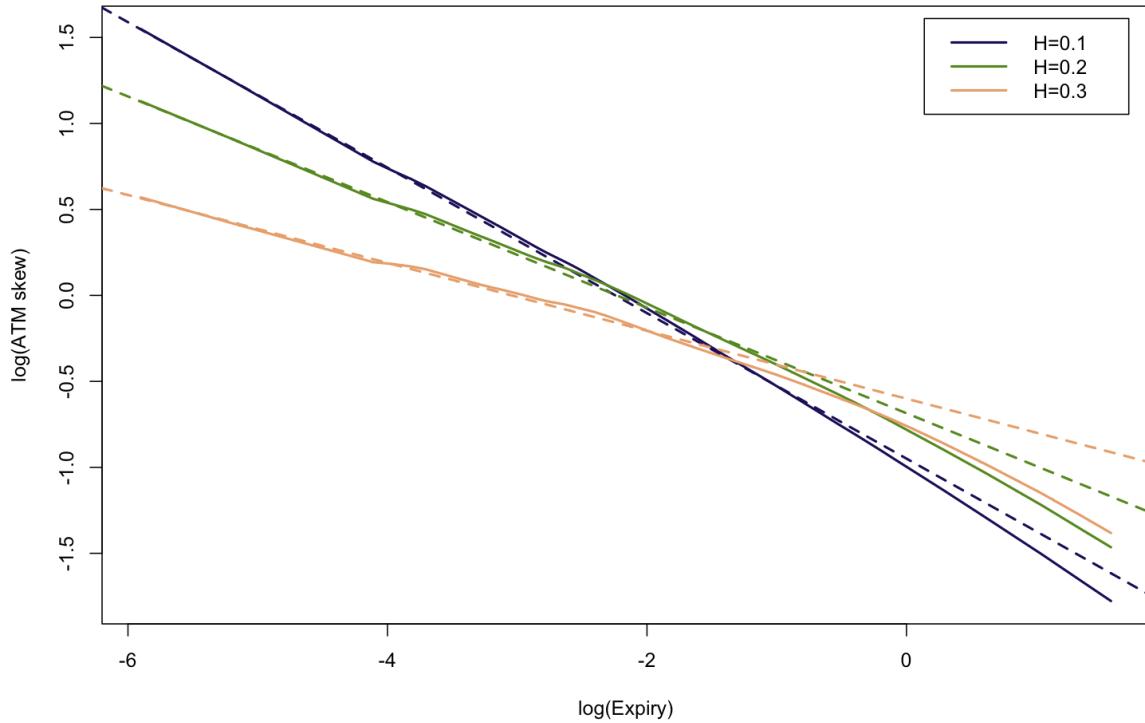


Figure 10:  $\log(\text{ATM skew})$  for  $H \in \{0.1, 0.2, 0.3\}$  together with linear fits to the first ten points.

### Estimate $H$ from term structure of skew

- We see that for short expirations, the rough Bergomi skew is almost a perfect power-law
  - $S(\tau) \sim \tau^{\alpha-1}$  with  $\alpha = H + 1/2$ .
- So let's estimate the slope of the empirical ATM skew!

### Estimate the power-law

```
In [24]: res.plot2013 <- plotIvolsMC(ivolsData, mcMatrix=mcMatrix2013, plot=F)

summary(fit.lm2013 <- lm(log(-res.plot2013$atmSkew[2:20])~log(res.plot2013$expiries[2:20])))

Call:
lm(formula = log(-res.plot2013$atmSkew[2:20]) ~ log(res.plot2013$expiries[2:20]))

Residuals:
    Min      1Q  Median      3Q     Max 
-0.039155 -0.014162 -0.008391  0.017412  0.045638 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) -0.979911   0.028417 -34.48 < 2e-16 ***
log(res.plot2013$expiries[2:20]) -0.240428   0.008568 -28.06 1.11e-15 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.02454 on 17 degrees of freedom
Multiple R-squared:  0.9789,    Adjusted R-squared:  0.9776 
F-statistic: 787.5 on 1 and 17 DF,  p-value: 1.113e-15
```

## Graph the fit

```
In [25]: plot(log(res.plot2013$expiries), log(-res.plot2013$atmSkew), pch=20, col=bl, xlab="Expiry", ylab="")  
abline(fit.lm2013, col=rd, lwd=2)
```

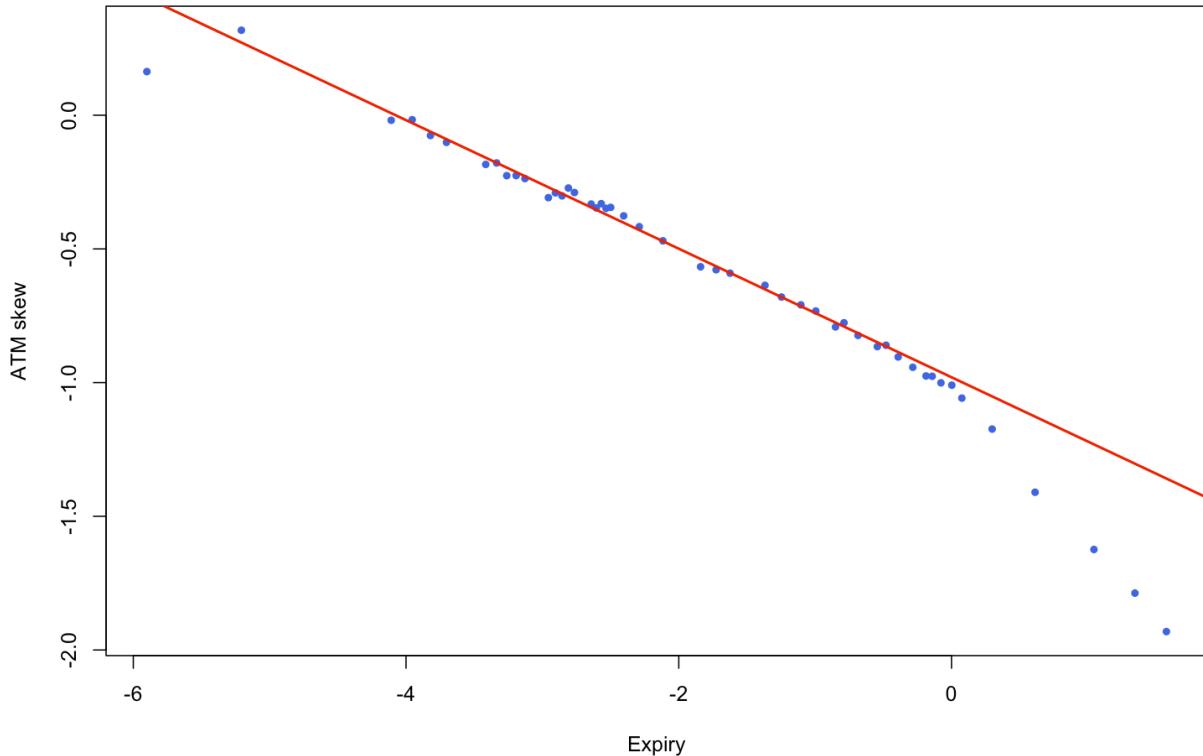


Figure 10: Blue points are empirical skews; the red line is a linear fit to the first 20 points (excluding the first).

## Run the rough Bergomi code with $H = 0.26$

- And playing with  $\eta$  and  $\rho$  a little...

```
In [26]: params.rBergomi.skew <- list(al=1-0.24, eta=1.8, rho=-0.8)  
xiCurve <- xi.curve.fast # We use the instantiated curve for speed  
  
t0<-proc.time()  
  
#number of iterations  
iters<- max(1,floor(paths/1000))  
  
#setup parallel backend  
cl.num <- detectCores() # This number is 8 on my MacBook Pro  
cl<-makeCluster(cl.num)  
registerDoParallel(cl)  
  
#loop  
ls <- foreach(icount(iters),.packages = "stinepack") %dopar% {  
  hybridSchemeRL.S(params.rBergomi.skew,xiCurve)(paths=1000, steps=steps, expiries=expi  
}  
  
stopCluster(cl)  
mcMatrix2023 <- do.call(cbind, ls) #Bind all of the submatrices into one big matrix
```

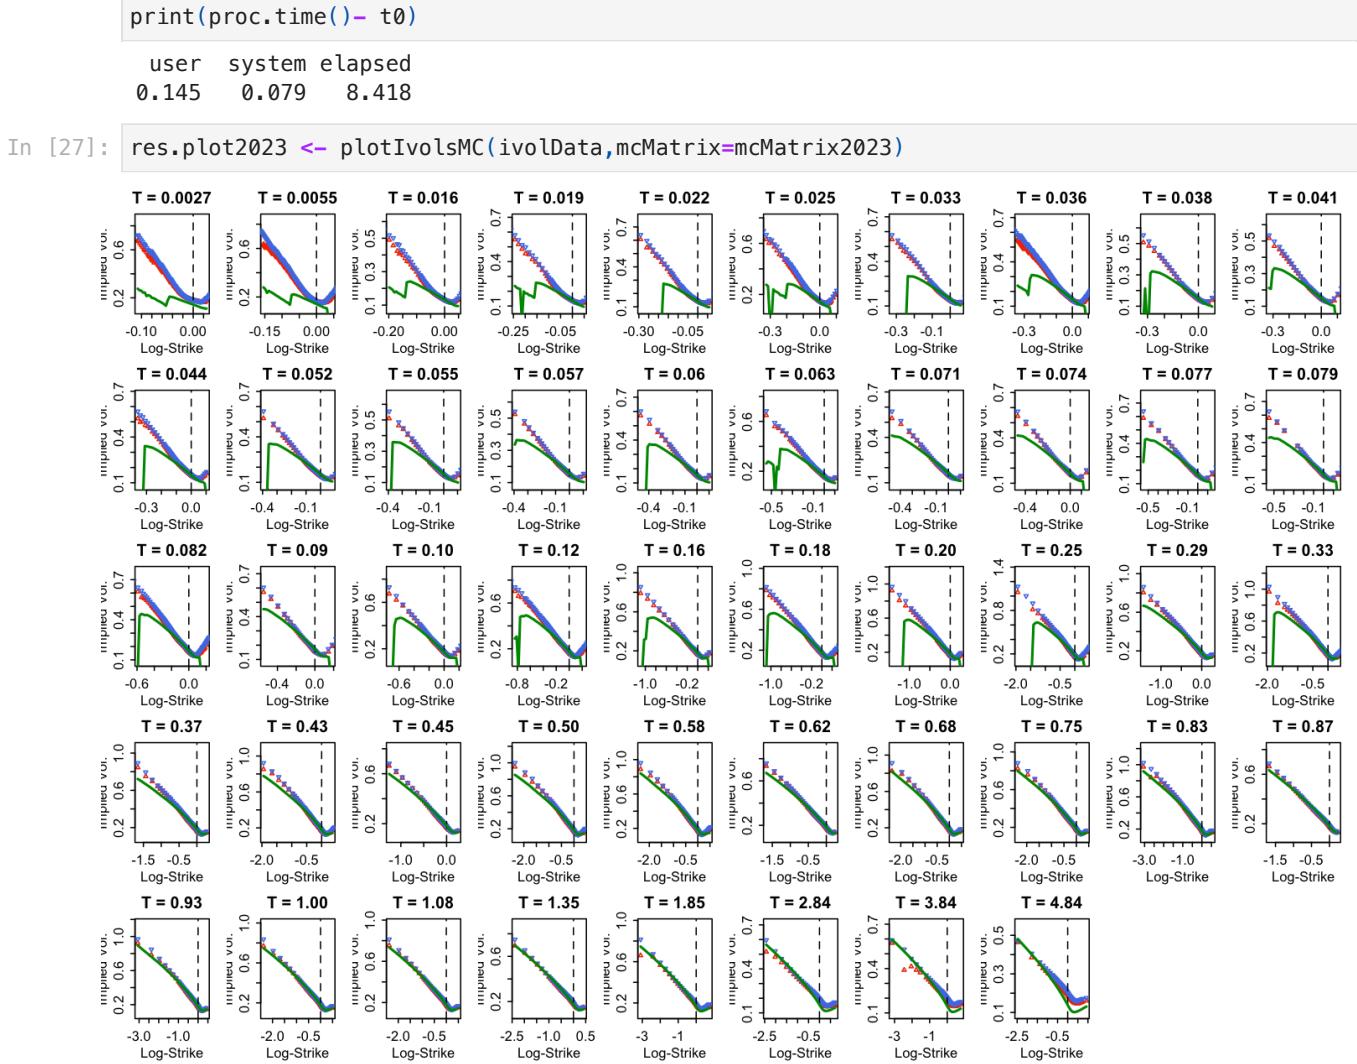


Figure 11: Blue points are empirical skews; the green lines are from the rough Bergomi simulation.

## Check the fit

```
In [28]: plot(log(res.plot2023$expiries), log(-res.plot2023$atmSkew), pch=20, col=bl, xlab="log(Expiry)", y
lines(log(res.plot2023$expiries), log(-res.plot2023$atmSkewMC), col=gr, lwd=2)
```

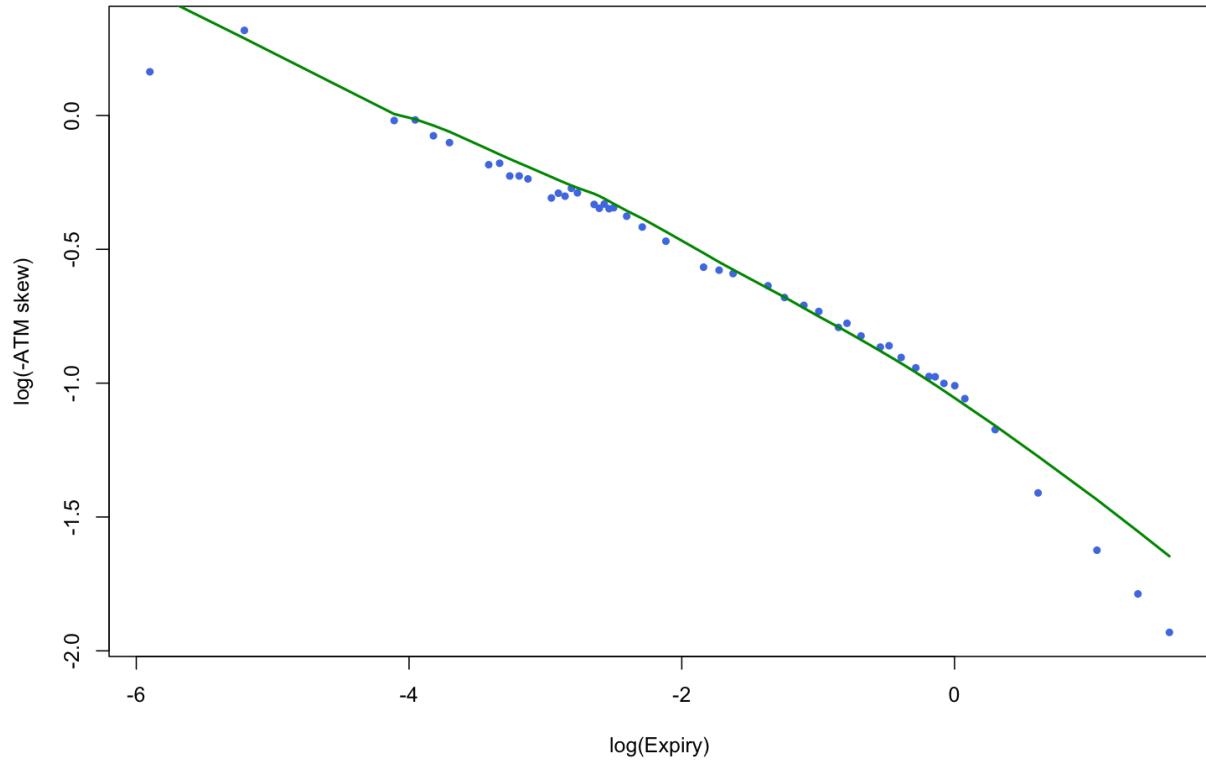


Figure 12: Blue points are empirical skews; the green line is from the rough Bergomi simulation.

### Check the six expirations, comparing with the 2013 guess

```
In [29]: res.plot2023.6 <- plotIvolsMC2(ivolData, mcMatrix=mcMatrix2023, mcMatrix2=mcMatrix2013, slices=
```

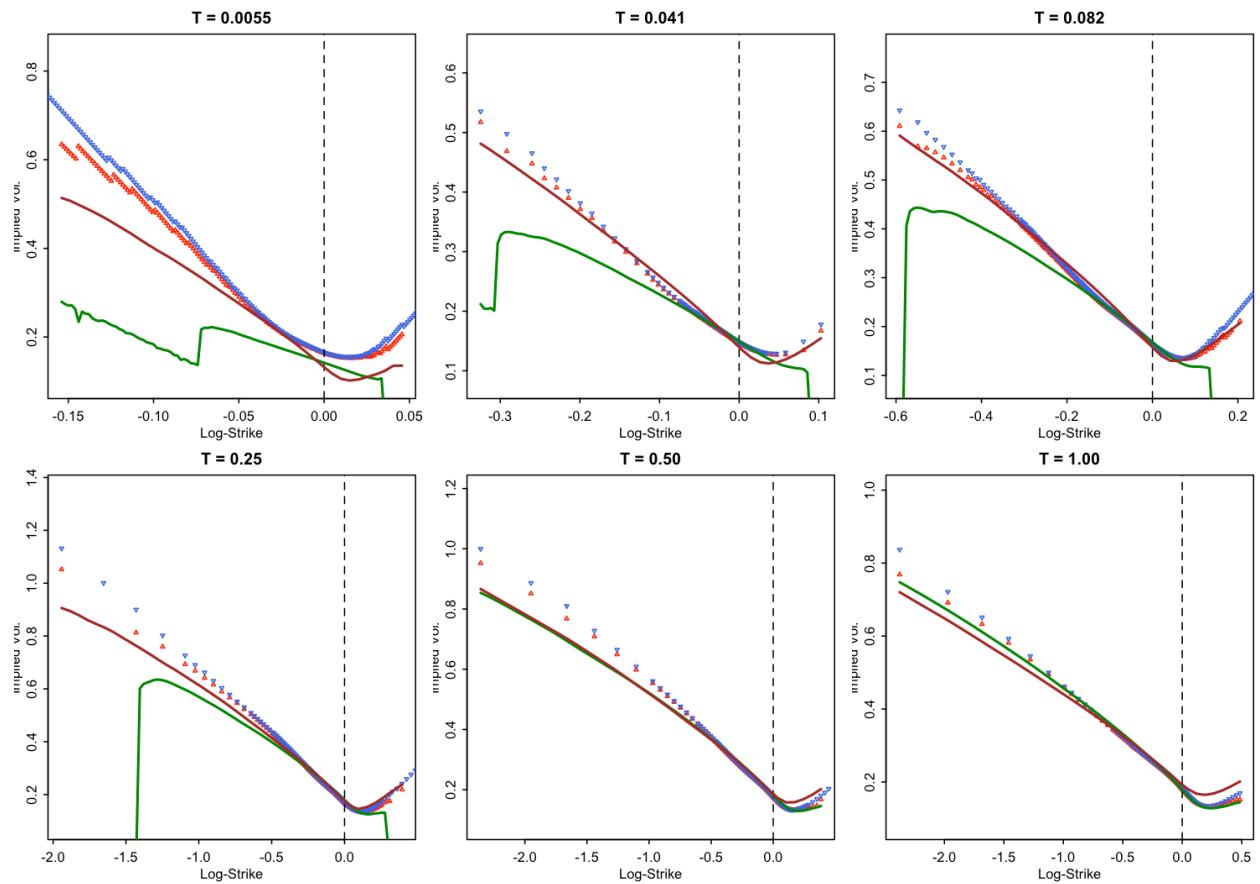


Figure 13: Six rough Bergomi smiles. Green is with parameters  $H = 0.26, \eta = 1.8, \rho = -0.8$ ; Brown is with parameters from 2013.

## Calibration using machine learning

- In a very well-cited paper, [Horvath et al.]<sup>[7]</sup> showed how to calibrate the rough Bergomi model to the volatility surface using machine learning.
  - A neural network is trained to map the shape of the volatility surface to model parameters.

## $H$ from VIX options and futures

- Rather than brute-force fitting a rough volatility model to the volatility surface, following [Jacquier, Martini and Muguruza], one can try to fix  $H$  from the term structure of the convexity adjustment between the variance contract and VIX futures.
- Once the Volterra process  $\tilde{W}$  has been simulated for this  $H$ , iterating on the parameters  $\eta$  and  $\rho$  to fit the observed volatility surface is relatively fast.

## The distribution of VIX future payoffs

- Denote the terminal value of the VIX futures by  $\sqrt{\zeta(T)}$ . Then, by definition (see Chapter 11 of [The Volatility Surface]<sup>[6]</sup> for more details),

$$\zeta(T) = \frac{1}{\Delta} \int_T^{T+\Delta} \mathbb{E}_T[V_u] du.$$

where  $\Delta$  is one month.

- In the rough Bergomi model,

$$V_u = \xi_t(u) \mathcal{E} \left( \eta \sqrt{2H} \int_t^u \frac{dW_s}{(u-s)^\gamma} \right)$$

with  $\gamma = 1/2 - H$  so  $V_u$  is lognormal.

## The lognormal approximation under rough Bergomi

- Under rough Bergomi, the VIX payoff and its square  $\zeta(T)$  should be approximately lognormally distributed.
  - The quality of this approximation was confirmed by [Jacquier, Martini and Muguruza]<sup>[8]</sup>.
  - In that case, the terminal distribution of  $\zeta(T)$  is completely determined by  $\mathbb{E}[\zeta(T)|\mathcal{F}_t]$  and  $\text{var}[\log \zeta(T)|\mathcal{F}_t]$ .
- Obviously

$$\mathbb{E}_t[\zeta(T)] = \frac{1}{\Delta} \int_T^{T+\Delta} \xi_t(u) du.$$

- Recall that forward variances  $\xi_t(u)$  may be estimated from variance swaps which can themselves be proxied by the log-strip (see Chapter 11 of [The Volatility Surface]<sup>[5]</sup> again).
  - Alternatively they may be estimated from linear strips of VIX options.

## Approximating the conditional variance of $\zeta(T)$ under rough Bergomi

- To estimate the conditional variance of  $\zeta(T)$ , we approximate the arithmetic mean by the geometric mean as follows:

$$\zeta(T) \approx \exp \left\{ \frac{1}{\Delta} \int_T^{T+\Delta} \mathbb{E}[\log V_u | \mathcal{F}_T] du \right\}.$$

Let  $y_u = \log V_u$  and recall that  $\gamma = \frac{1}{2} - H$ . Apart from  $\mathcal{F}_t$  measurable terms (abbreviated as ``drift''), we have

$$\begin{aligned} \int_T^{T+\Delta} \mathbb{E}_t[y_u] du &= \eta \sqrt{2H} \int_t^T \frac{dW_s}{(u-s)^\gamma} du + \text{drift} \\ &= \eta \sqrt{2H} \int_t^T \int_T^{T+\Delta} \frac{du}{(u-s)^\gamma} dW_s + \text{drift} \\ &= \eta \frac{\sqrt{2H}}{1-\gamma} \int_t^T [(T+\Delta-s)^{1-\gamma} - (T-s)^{1-\gamma}] dW_s + \text{drift}. \end{aligned}$$

This gives

$$\begin{aligned} \text{var}[\log \zeta(T) | \mathcal{F}_t] &\approx \frac{\eta^2}{\Delta^2} \frac{2H}{(H+1/2)^2} \int_t^T [(T+\Delta-s)^{1/2+H} - (T-s)^{1/2+H}]^2 ds \\ &= \eta^2 (T-t)^{2H} f^H \left( \frac{\Delta}{T-t} \right) \end{aligned}$$

where

$$f^H(\theta) = \frac{2H}{(H+1/2)^2} \frac{1}{\theta^2} \int_0^1 \left[ (1+\theta-x)^{1/2+H} - (1-x)^{1/2+H} \right]^2 dx.$$

## The approximate fair value of VIX futures

- In [Bayer, Friz and Gatheral]<sup>[1]</sup>, we chose to study the term structure of VVIX (the VIX of VIX).
- It is more natural to follow [Jacquier, Martini and Muguruza]<sup>[9]</sup> and approximate the fair value of VIX futures.
- Under the lognormal approximation, the fair value of the  $T$ -maturity VIX future is given by

$$\mathbb{E} \left[ \sqrt{\zeta(T)} \mid \mathcal{F}_t \right] = \sqrt{\mathbb{E} [\zeta(T) \mid \mathcal{F}_t]} \exp \left\{ -\frac{1}{8} \text{var}[\log \zeta(T) \mid \mathcal{F}_t] \right\}.$$

## Load VIX option data

```
In [30]: # load("vixIvols20230215.rData")
vixVolData <- vixIvols20230215
head(vixVolData)
```

A data.frame: 6 × 7

	Expiry	Texp	Strike	Bid	Ask	Fwd	CallMid
	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	20230222	0.01916496	10.0	NA	2.086124	20.19517	NA
2	20230222	0.01916496	10.5	NA	1.952121	20.19517	NA
3	20230222	0.01916496	11.0	NA	1.995202	20.19517	NA
4	20230222	0.01916496	11.5	NA	1.862784	20.19517	NA
5	20230222	0.01916496	12.0	NA	1.735648	20.19517	NA
6	20230222	0.01916496	12.5	NA	1.613290	20.19517	NA

```
In [31]: plotVIX <- plotIvols(vixVolData)
```

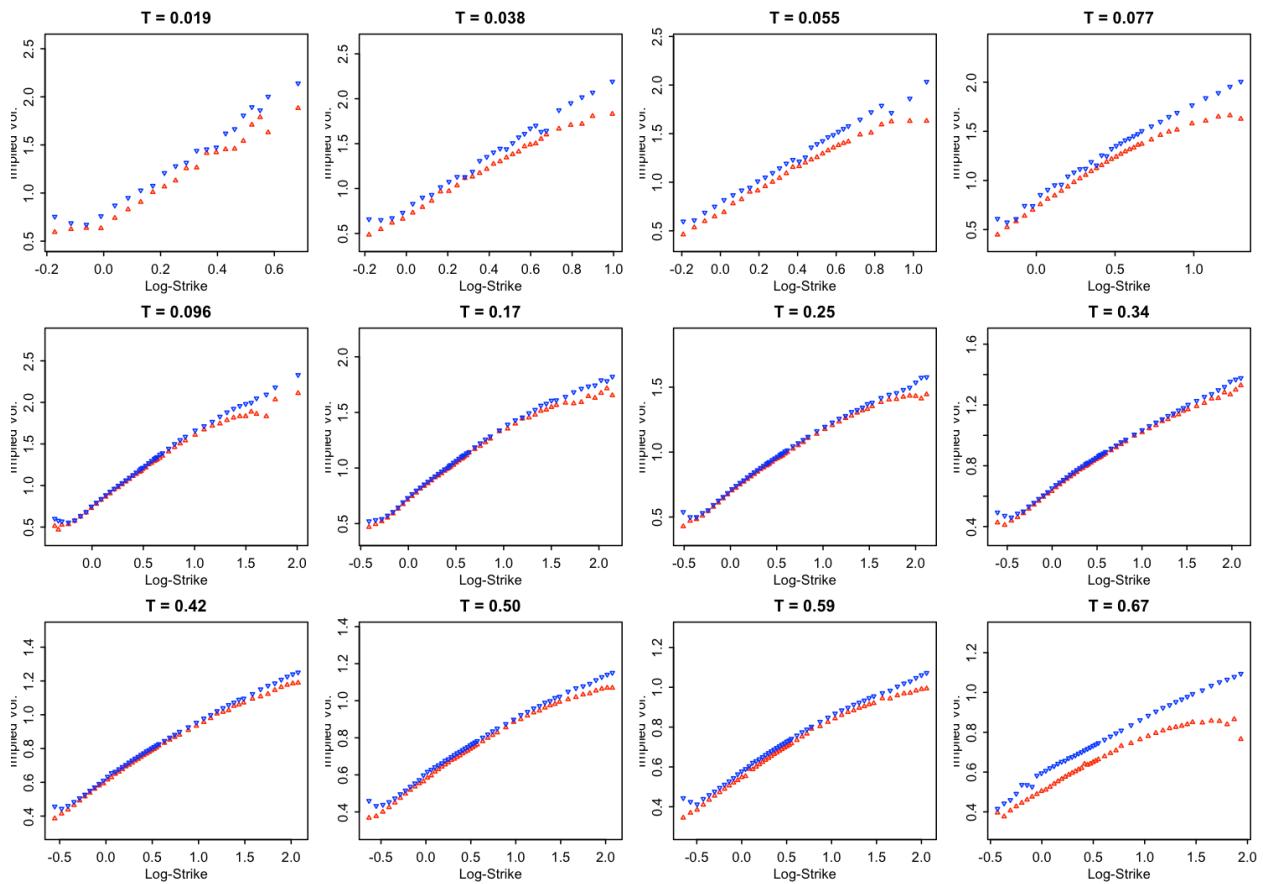


Figure 14: VIX smiles as of February 15, 2023.

## VIX futures from put-call parity

```
In [32]: t.VIX <- plotVIX$expiries
(f.VIX <- unique(vixVolData$Fwd))

20.1951741855249 · 20.3550062990049 · 20.6097782540332 · 20.468796922825 · 20.1261648745547 ·
21.1644803229044 · 21.6400670851342 · 22.0432737879408 · 22.5994173344876 · 22.608388198057 ·
23.0021994712784 · 23.0946832691182
```

```
In [33]: plot(t.VIX,f.VIX,pch=20,col=bl,cex=2,type="b")
```

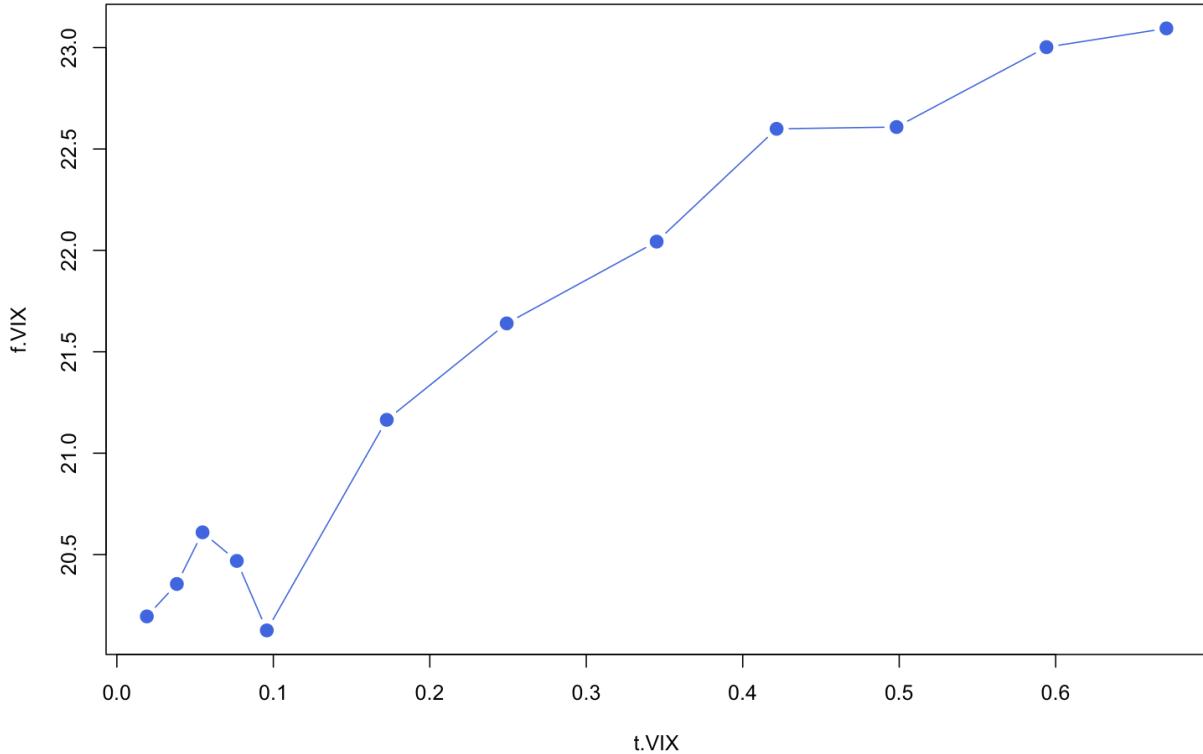


Figure 15: The VIX futures curve from put-call parity

### $\mathbb{E}_t [\zeta(T)]$ from VIX option data

- We can span the payoff of a forward starting variance swap  $\zeta(T) = \mathbb{E}_T \left[ \int_T^{T+\Delta} V_s ds \right]$  using VIX options.
- From the Carr-Madan spanning formula,

$$\mathbb{E}_t [\zeta(T)] = F_{VIX}^2 + 2 \int_0^{F_{VIX}} P(K) dK + 2 \int_{F_{VIX}}^{\infty} C(K) dK.$$

- We need to interpolate and extrapolate out-of-the-money option prices to get the *convexity adjustment*.

### Interpolation and extrapolation of VIX smiles

- We perform
  - Monotonic spline interpolation of mid-vols.
  - Extrapolation at constant level.
- Then we integrate the Black-Scholes formula with these vols.

```
In [34]: vix2 <- function(ivolData) function(slice){
  bidVols <- as.numeric(ivolData$Bid)
  askVols <- as.numeric(ivolData$Ask)
  expDates <- unique(ivolData$Texp)
```

```
#####
# Interpolate and extrapolate vols for this slice at requested output points

t <- expDates[slice]
texp <- ivolData$Texp
bidVol <- bidVols[texp==t]
askVol <- askVols[texp==t]
midVol <- (bidVol+askVol)/2
f <- (ivolData$Fwd[texp==t])[1]
k <- log(ivolData$Strike[texp==t]/f) # Plot vs log-strike
include <- !is.na(bidVol)
kmin <- min(k[include])
kmax <- max(k[include])

# Compute and store interpolated and extrapolated vols
kIn <- k[!is.na(midVol)]
volIn <- midVol[!is.na(midVol)]
volInterp <- function(kout){
  if (kout < kmin){res <- midVol[which(k==kmin)] }
  else if( kout > kmax){res <- midVol[which(k==kmax)] }
  else res <- stinterp(x=kIn,y=volIn, kout)$y
  return(res)
}
vixVol <- function(x){sapply(x,volInterp)}
# Now we use the vectorized function vixVol to compute the convexity adjustemnt
cTilde <- function(y){exp(y)*BSFormula(1, exp(y), t, r=0, vixVol(y))}
pTilde <- function(y){exp(y)*BSFormulaPut(1, exp(y), t, r=0, vixVol(y))}
callIntegral <- integrate(cTilde,lower=0,upper=10)$value
putIntegral <- integrate(pTilde,lower=-10,upper=0)$value
res <- f^2*(1+2*(callIntegral+putIntegral))
return(res)
}
```

In [35]: `(e.VIX2 <- Vectorize(vix2(vixVolData))(1:12)/10^4)`

```
0.0414436421593004 · 0.0429423797669813 · 0.0447453688499278 · 0.0455250592633412 ·  
0.0454742634194872 · 0.0540602065324636 · 0.0593273396600168 · 0.0639020315523995 ·  
0.0687341512776171 · 0.0698957062653717 · 0.0744212383905677 · 0.0767311559552205
```

## Functions to compute the convexity adjustment

In [36]:

```
etaNu <- function(nu,h){
  ch2 <- gamma(3/2-h)/(gamma(h+1/2)*gamma(2-2*h))
  return(2*nu*sqrt(ch2))
}

fH <- function(theta,h){
  integ <- function(x){((1+theta-x)^(h+1/2)-(1-x)^(h+1/2))^2}
  tmp <- integrate(integ, lower=0, upper=1)$value
  return(tmp/theta^2*(2*h)/(h+1/2)^2)
}

varLogPsi <- function(tau, delta, h, nu){etaNu(nu,h)^2*tau^(2*h)*fH(delta/tau,h)}

convAdj.raw <- function(tau, delta, h, nu){exp(-varLogPsi(tau, delta, h, nu)/8)}
convAdj <- Vectorize(convAdj.raw, vectorize.args="tau")
```

## Function to compute VIX futures given $H$ and $\nu$

In [37]:

```
e.VIX <- function(paramvec) function(t.VIX,e.VIX2){

  H <- paramvec[1]
  eta <- paramvec[2]
```

```

    conv.adj <- convAdj(t.VIX, delta=1/12, H, eta)
    return(sqrt(e.VIX2) * conv.adj * 100)
}

```

## Fix dates for optimization

```
In [38]: eVIX <- function(paramvec){

  e.VIX <- e.VIX(paramvec)(t.VIX,e.VIX2)
  return(e.VIX)

}

eVIX(c(.2,.6))
```

20.2091390147171 · 20.457329893908 · 20.7977286657498 · 20.8784895162704 · 20.7892775905226 ·  
 22.389710052101 · 23.2257302822626 · 23.8581403960556 · 24.5673339020038 · 24.6142080585121 ·  
 25.2124237273954 · 25.4624172549922

Compare with actual futures curve.

```
In [39]: f.VIX

20.1951741855249 · 20.3550062990049 · 20.6097782540332 · 20.468796922825 · 20.1261648745547 ·  

21.1644803229044 · 21.6400670851342 · 22.0432737879408 · 22.5994173344876 · 22.608388198057 ·  

23.0021994712784 · 23.0946832691182
```

## Optimize

```
In [40]: obj <- function(paramvec){

  eVIX.model <- eVIX(paramvec)
  return(sum((eVIX.model-f.VIX)^2)*1e6)
}
```

```
In [41]: (res.20230215.VIX <- optim(c(.3,.15),obj,method="L-BFGS-B",lower=c(0.0001,0.01),upper=c(.49,10))

$par          0.185399648790267 · 0.916948255714096
$value        337483.626861749
$counts       function: 28 gradient: 28
$convergence  0
$message      'CONVERGENCE: REL_REDUCTION_OF_F <= FACTR*EPSMCH'
```

## Plot actual and fitted VIX futures

```
In [42]: plot(t.VIX,f.VIX,pch=20,col=bl,cex=2,ylim=c(20,23.5))
points(t.VIX,eVIX(res.20230215.VIX$par),pch=20,col=pk,cex=1)
```

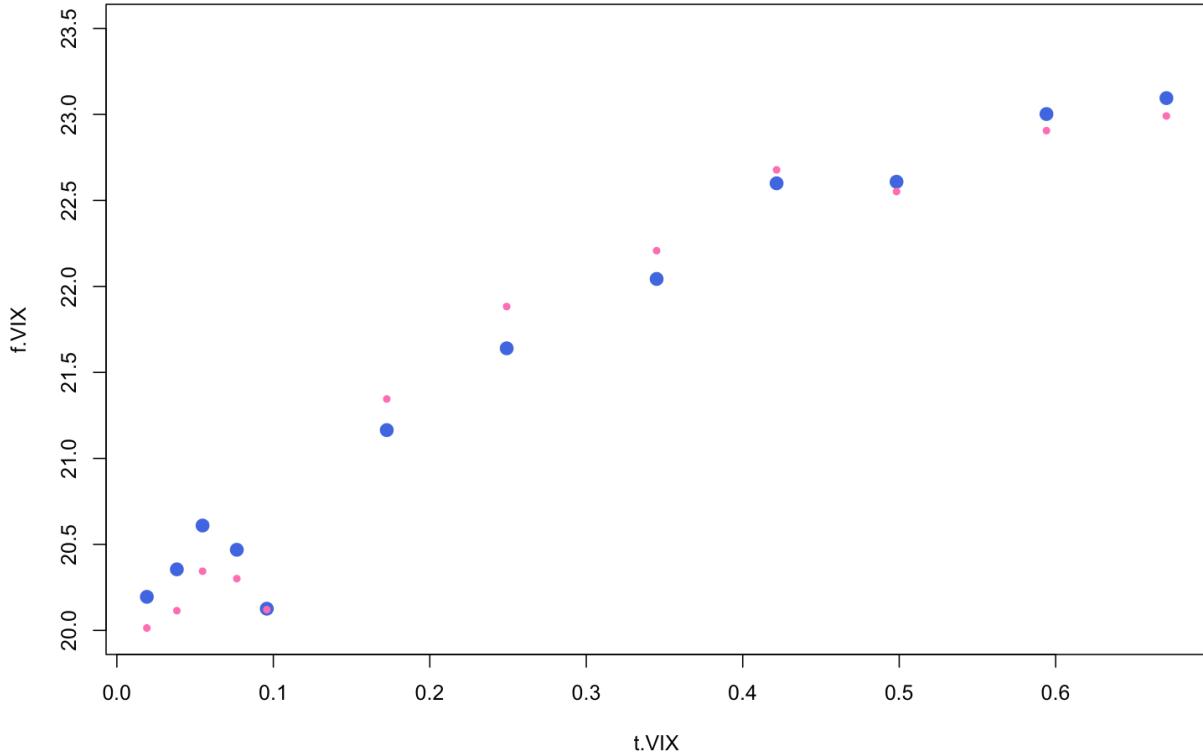


Figure 16: Fit of the rough Bergomi VIX approximation to the observed VIX futures curve. VIX formula in blue; actual VIX futures in pink.

## Summary

- The VIX estimate gives  $H = 0.185$ .

Now draw smiles with  $H = 0.185$

```
In [43]: params.rBergomi.VIX <- list(al=0.685, eta=1.7, rho=-0.8)
xiCurve <- xi.curve.fast # We use the instantiated curve for speed

In [44]: t0<-proc.time()

#number of iterations
iters<- max(1,floor(paths/1000))

#setup parallel backend
cl.num <- detectCores() # This number is 8 on my MacBook Pro
cl<-makeCluster(cl.num)
registerDoParallel(cl)

#loop
ls <- foreach(icount(iters),.packages = "stinepack") %dopar% {
  hybridSchemeRL.S(params.rBergomi.VIX,xiCurve)(paths=1000, steps=steps, expiries=expir)

stopCluster(cl)
mcMatrix.VIX <- do.call(cbind, ls) #Bind all of the submatrices into one big matrix

print(proc.time()- t0)
```

user	system	elapsed
0.144	0.068	8.432

## Plot actual and rough Bergomi smiles

```
In [45]: res.plot.VIX <- plotIvolsMC(ivolData, mcMatrix=mcMatrix.VIX, plot=F)

In [46]: plot(res.plot.VIX$expiries, res.plot.VIX$atmSkew, pch=20, col=bl, xlab="Expiry", ylab="ATM skew")
          lines(res.plot.VIX$expiries, res.plot.VIX$atmSkewMC, col=gr, lwd=2)
```

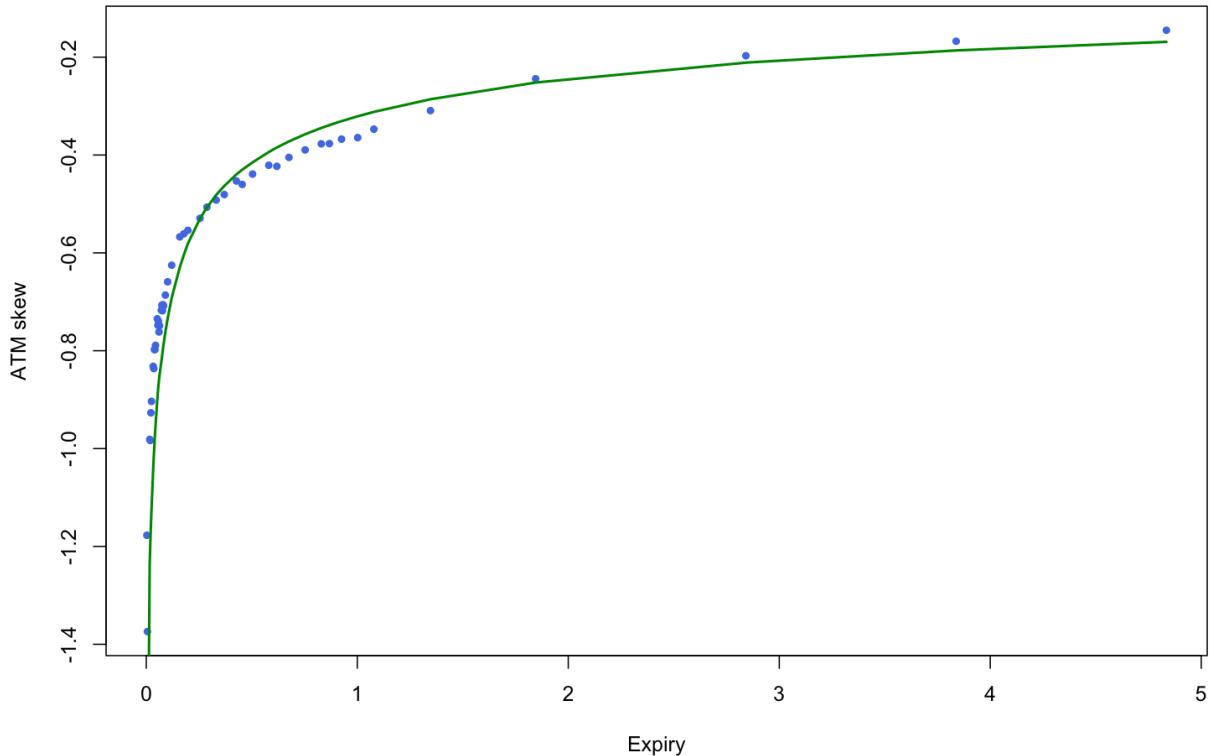


Figure 17: Actual (in blue) vs fitted (in green) SPX ATM skew.

## Compare smiles with the two choices of $H$

```
In [47]: res.plot4 <- plotIvolsMC2(ivolData, mcMatrix=mcMatrix.VIX, mcMatrix2=mcMatrix2023,, slices= c(2,,
```

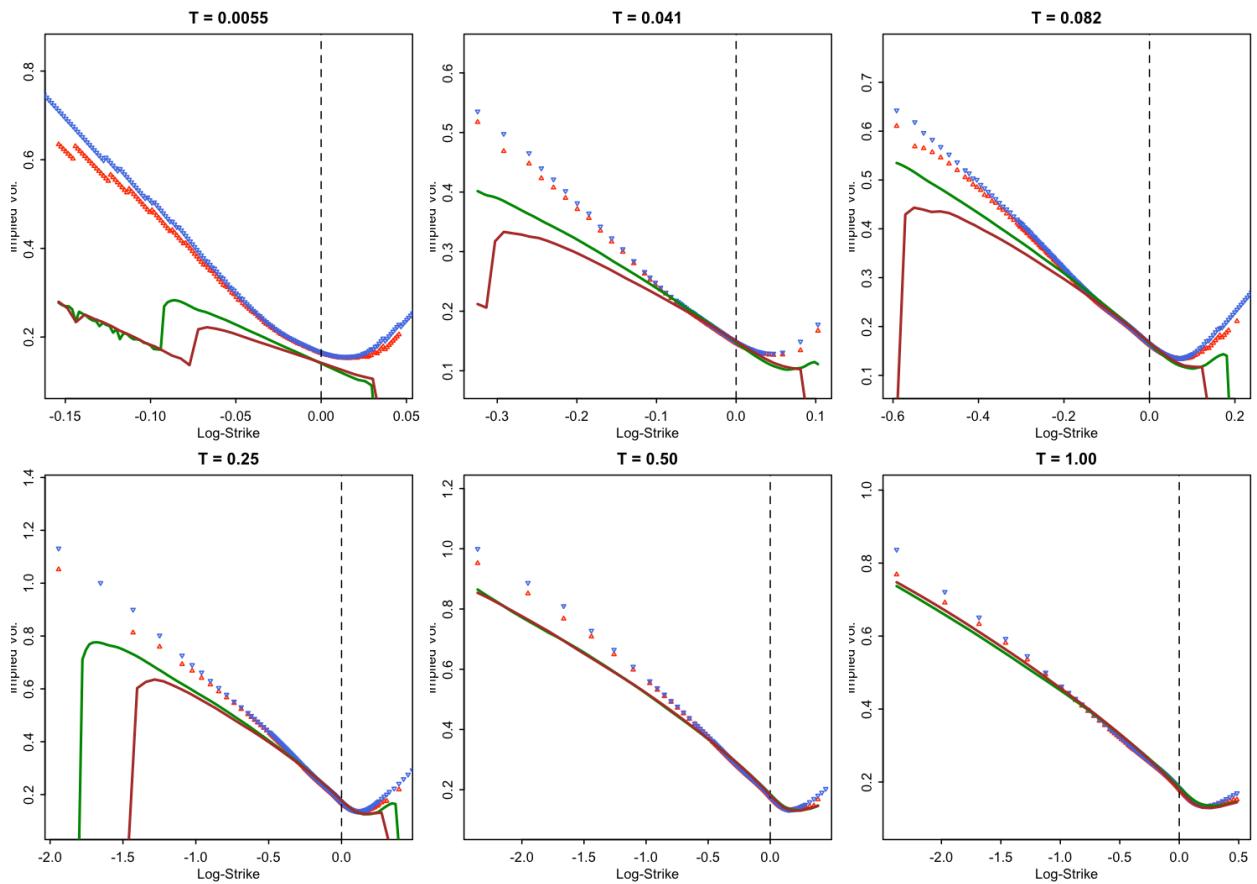


Figure 18: Six rough Bergomi smiles. Green is with parameters `params.rBergomi.VIX`; brown s with parameters `params.rBergomi.skew`.

## Rough Bergomi parameters under $\mathbb{P}$ and under $\mathbb{Q}$

- We might wonder whether implied model parameters are consistent with historical parameters.
- It is shown in [Bayer, Friz and Gatheral]<sup>[2]</sup> that the volatility of volatility parameter  $\eta$  in the rough Bergomi model and the volatility of volatility  $\nu$  in the historical time series should be related as follows.

$$\tilde{\eta} := \eta \sqrt{2H} = 2\nu C_H$$

with

$$C_H = \sqrt{\frac{2H\Gamma(3/2-H)}{\Gamma(H+1/2)\Gamma(2-2H)}}.$$

## Parameter estimates under $\mathbb{Q}$

In Section 5.2 of [Bayer, Friz and Gatheral]<sup>[1]</sup>, parameter guesses for the SPX implied volatility surface on two particular dates in history are given as follows:

Date	$H$	$\eta$	$\tilde{\eta}$
February 4, 2010	0.07	1.9	0.7109
August 14, 2013	0.05	2.3	0.7273

- Estimates of  $\tilde{\eta}$  seem more stable than estimates of  $\eta$  and  $H$  separately.
- We observe the same phenomenon when estimating  $\nu$  and  $H$  from historical RV data.
  - Estimates of the product  $\nu \sqrt{H}$  are more stable than estimates of the two parameters separately.

## Parameter estimates under $\mathbb{P}$

- From our analysis of the SPX realized variance time series in Lecture 1, we estimated

$$H \approx 0.166, \quad \nu \approx 0.302.$$

- Plugging these estimates into the formula (from above)

$$\tilde{\eta}_1 = 2\nu \sqrt{\frac{2H\Gamma(3/2-H)}{\Gamma(H+1/2)\Gamma(2-2H)}} \approx 0.268.$$

```
In [48]: h.est <- 0.166
nu.est <- 0.302
(nu.tilde <- 2*nu.est*sqrt(2*h.est*gamma(3/2-h.est)/gamma(h.est+1/2)*gamma(2-2*h.est)))
```

0.268425542130751

- Seemingly inconsistent with the implied estimate of around 0.72.

```
In [49]: nu.tilde*252^h.est
```

0.672134717914495

- However, the historical estimate is in daily terms and the implied estimate in annualized terms.
- To convert, we need to multiply the historical estimate by the annualization factor  $(252)^H$ , to get

$$\tilde{\eta} \approx \tilde{\eta}_1 \times (252)^H = 0.67.$$

- Historical and implied estimates are consistent.

## Rough volatility and long memory

- In [Bennedsen, Lunde and Pakkanen]<sup>[4]</sup>, the authors show how we can both have our cake and eat it by choosing different kernels.
- In particular, with appropriate choices of  $\gamma$  and  $\beta$  the kernel

$$\kappa(\tau) = \frac{1}{\tau^\gamma (1+\tau)^\beta}$$

generates a model that exhibits both rough volatility and power-law decay of the autocorrelation function.  
 - That is rough volatility plus long memory.

- Models with more parameters may of course also fit the volatility surface better.

## Forecasting the variance swap curve

In [Bayer, Friz and Gatheral]<sup>[1]</sup>, we show how to forecast the whole variance swap curve using the variance forecasting formula.

- We show consistency between the volatility forecast under  $\mathbb{P}$  and the forward variance curve (under  $\mathbb{Q}$ ) around two of the most dramatic events:
  - The collapse of Lehman Brothers, and
  - The Flash Crash.

## Features of the rough Bergomi model

- In Lecture 1, scaling properties of the time series of historical volatility suggested a natural non-Markovian stochastic volatility model under  $\mathbb{P}$ .
- The simplest specification of  $\frac{d\mathbb{Q}}{d\mathbb{P}}$  gives the rough Bergomi model, a non-Markovian generalization of the Bergomi model.
  - The history of the Brownian motion  $\{W_s, s < t\}$  required for pricing is encoded in the forward variance curve, which is observed in the market.
  - Efficient computations are possible using the hybrid BSS scheme.
- Rough Bergomi is easy to simulate using the hybrid-BSS scheme.
- Rough Bergomi is a lognormal model and thus has reasonable dynamics.
- However, rough Bergomi gives flat VIX smiles.

## More rough volatility models

This form suggests many other rough volatility models of the form

$$\begin{aligned}\frac{dS_t}{S_t} &= \sqrt{\xi_t(t)} dZ_t \\ d\xi_t(u) &= \lambda(\xi) \kappa(u - t) dW_t\end{aligned}$$

where both the function  $\lambda$  and the kernel  $\kappa$  depend on the model.

- As long as  $\kappa(\tau) \sim \tau^{-\gamma}$  as  $\tau \rightarrow 0$ , the model will be rough in the sense that sample paths of instantaneous variance will be Hölder continuous with exponent  $H = \frac{1}{2} - \gamma$ .

## The rough Heston model

By considering the limit of a simple Hawkes process-based model of order flow, [Jaisson and Rosenbaum]<sup>[9]</sup> and [El Euch, Fukasawa and Rosenbaum]<sup>[5]</sup> derive a rough Heston model. The equation for variance in this model takes the form

$$V_u = \theta_t(u) - \frac{1}{\Gamma(\alpha)} \int_t^u (u-s)^{\alpha-1} \lambda V_s ds + \frac{1}{\Gamma(\alpha)} \int_t^u (u-s)^{\alpha-1} \nu \sqrt{V_s} dW_s.$$

where  $\alpha = H + \frac{1}{2}$ .

- $H \in (0, \frac{1}{2}]$  is the Hurst exponent of the volatility,  $\lambda > 0$  is the mean reversion parameter,  $\eta > 0$  is the volatility of volatility parameter.
- The function  $\theta$  is assumed to be continuous and represents a time-dependent mean reversion level.

- The rough Heston model generalizes the classical Heston model which is recovered when  $H = 1/2$ .

## Forward variance in the rough Heston model ( $\lambda = 0$ )

- We will consider only the special case  $\lambda = 0$ . In this case,  $\xi_t(u) = \mathbb{E}_t[V_u] = \theta_t(u)$ .
- It follows that

$$V_u = \xi_t(u) + \frac{\nu}{\Gamma(\alpha)} \int_t^u (u-s)^{\alpha-1} \sqrt{V_s} dW_s.$$

- Also

$$V_u = \xi_{t+h}(u) + \frac{\nu}{\Gamma(\alpha)} \int_{t+h}^u (u-s)^{\alpha-1} \sqrt{V_s} dW_s.$$

## The rough Heston model with $\lambda = 0$ in forward variance form

Subtracting these two equations gives

$$\xi_{t+h}(u) - \xi_t(u) = \frac{\nu}{\Gamma(\alpha)} \int_t^{t+h} (u-s)^{\alpha-1} \sqrt{V_s} dW_s.$$

Taking the limit  $h \rightarrow 0$ , we obtain

$$d\xi_t(u) = \frac{\nu}{\Gamma(\alpha)} (u-t)^{\alpha-1} \sqrt{V_t} dW_t,$$

the rough Heston model in forward variance form.

## The rough Heston model in forward variance form ( $\lambda \geq 0$ )

Let  $\phi(\tau) = \frac{1}{\Gamma(\alpha)} \tau^{\alpha-1}$  and  $\kappa(\tau) = \tau^{\alpha-1} E_{\alpha,\alpha}(-\lambda\tau^\alpha)$ . Then

$$\phi - \kappa = \lambda (\phi \star \kappa).$$

That is,  $\phi$  is the  $\lambda$ -resolvent of  $\kappa$ . To check this, take Laplace transforms:

$$\frac{1}{p^\alpha} - \frac{1}{p^\alpha + \lambda} = \lambda \frac{1}{p^\alpha} \frac{1}{p^\alpha + \lambda}.$$

## The rough Heston model in forward variance form ( $\lambda \geq 0$ )

Consider once again the rough Heston model of El Euch and Rosenbaum:

$$V_u = \theta_t(u) - \frac{\lambda}{\Gamma(\alpha)} \int_t^u (u-s)^{\alpha-1} V_s ds + \frac{1}{\Gamma(\alpha)} \int_t^u (u-s)^{\alpha-1} \nu \sqrt{V_s} dW_s.$$

Write this formally as:

$$V = \theta - \lambda (\phi \star V) + \nu (\phi \star \sqrt{V} dW).$$

Convolve with  $\kappa$  to get

$$(\kappa \star V) = (\kappa \star \theta) - \lambda (\kappa \star \phi \star V) + \nu (\kappa \star \phi \star \sqrt{V} dW).$$

But  $\phi - \kappa = \lambda (\phi \star \kappa)$  so

$$(\kappa \star V) = (\kappa \star \theta) - ((\phi - \kappa) \star V) + \frac{\nu}{\lambda} ((\phi - \kappa) \star \sqrt{V} dW).$$

Thus

$$0 = -\lambda (\kappa \star \theta) + \lambda (\phi \star V) - \nu (\phi \star \sqrt{V} dW) + (\kappa \star \sqrt{V} dW).$$

Add to the first equation to get

$$V = \theta - \lambda (\kappa \star \theta) + \nu (\kappa \star \sqrt{V} dW).$$

Thus, the drift term is eliminated and we must have  $\xi = \theta - \lambda (\kappa \star \theta)$ .

## The rough Heston model in forward variance form ( $\lambda \geq 0$ )

Writing out the last equation in full,

$$V_u = \xi_t(u) + \nu \int_t^u \kappa(u-s) \sqrt{V_s} dW_s.$$

With the same argument as before, we get the rough Heston model in forward variance form:

$$d\xi_t(u) = \nu \kappa(u-t) \sqrt{V_t} dW_t,$$

with  $\kappa(\tau) = \tau^{\alpha-1} E_{\alpha,\alpha}(-\lambda\tau^\alpha)$ .

## Non-Markovianity of the rough Heston model

- Note that the limit  $u \rightarrow t$  of the rough Heston model makes no sense.
  - This reflects the fact that the rough Heston model is not Markovian.
    - There is no SDE for  $V_t$  and no corresponding PDE.
  - On the other hand, we can write an SDE for each  $\xi_t(u)$ ,  $u > t$ .
    - We can even apply Itô's Lemma!
- The rough Heston model is Markovian in the infinite-dimensional forward variance curve  $\xi_t(u)$ ,  $u > t$ .

## Features of the rough Heston model

- The rough Heston model, as we will see in Lecture 3, is very tractable, at least with  $\lambda = 0$ .
  - Arguably more tractable than the classical Heston model.
- The rough Heston model arises as a limit of a simple Hawkes process-based model of order flow.
- The model is harder to simulate than rough Bergomi, but not so hard (see Lecture 4).
- However, since the model is affine, dynamics are not consistent with observation.
  - VIX smiles are negatively sloped, totally inconsistent with observation!

## Dynamics of the volatility surface: Model dependence

- All rough stochastic volatility models have essentially the same implications for the shape of the volatility surface.
- At first it might therefore seem that it would be hard to differentiate between models.

- That would certainly be the case if we were to confine our attention to the shape of the volatility surface today.
- If instead we were to study the dynamics of the volatility skew – in particular, how the observed volatility skew depends on the overall level of volatility, we would be able to differentiate between models.
- As explained in [The Volatility Surface]<sup>[9]</sup>, we expect the ATM volatility skew to be roughly independent of the ATM volatility in a lognormal model such as rough Bergomi.
- In Figure 4, we see how the ATM skew varies with ATM volatility under rough Bergomi and rough Heston, and compare with empirical estimates.

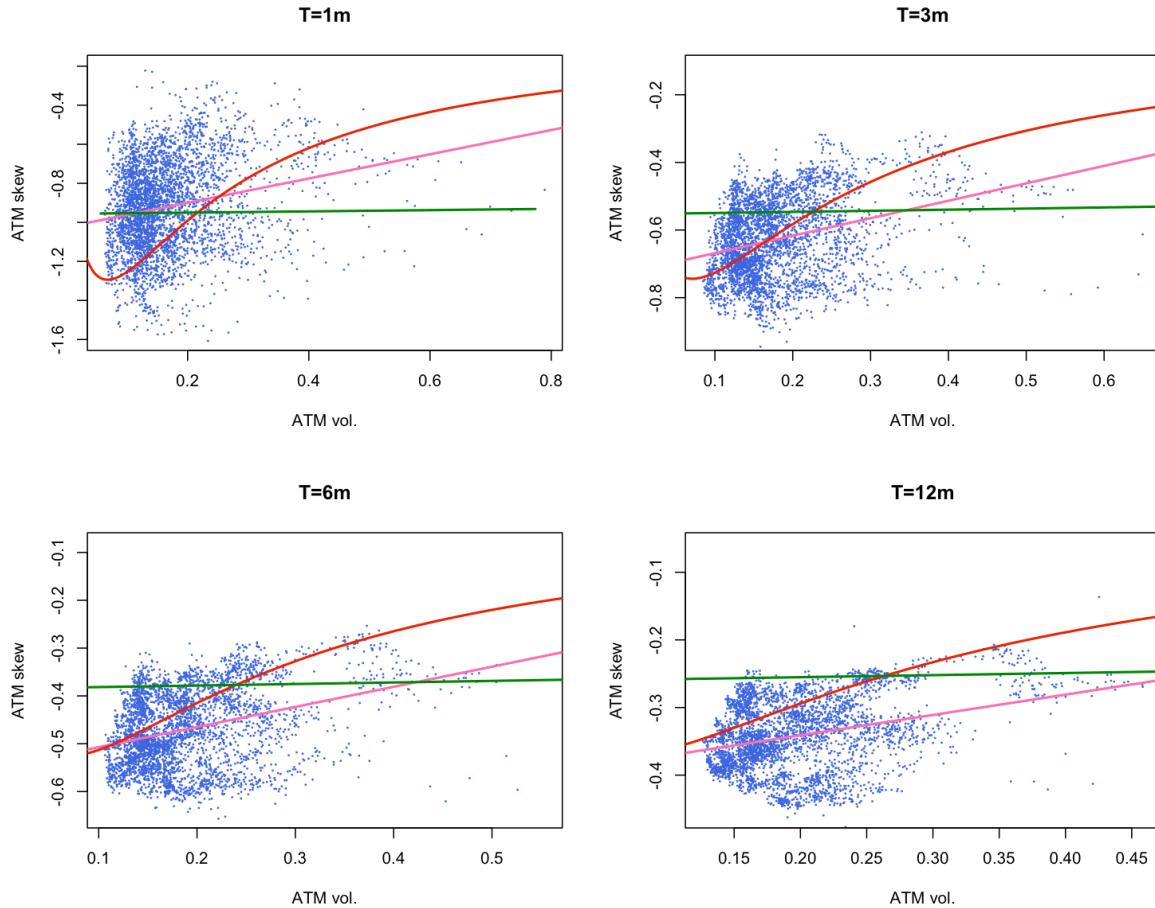


Figure 19: Blue points are empirical 3-month ATM volatilities and skews (from Sep-2008 to today); a regression line in pink; the green line is the rough Bergomi computation with the above parameters; the red line is rough Heston.

## The quadratic rough Heston model

The quadratic rough Heston (QRH) model of [QRH]<sup>[9]</sup> may be written as

$$\begin{aligned} \frac{dS_t}{S_t} &= -\sqrt{V_t} dW_t, \\ V_t &= Y_t^2 + c, \end{aligned} \tag{2}$$

where  $c \geq 0$ , and

$$Y_t = \bar{Y} + \int_{-\infty}^t \kappa(t-s) \sqrt{V_s} dW_s = \bar{Y} - \int_{-\infty}^t \kappa(t-s) \frac{dS_s}{S_s} \quad (3)$$

is a weighted average of historical returns and  $\kappa(\cdot)$  a kernel function.

## Dynamics in martingale form

Taking the conditional expectation of  $Y_u$  for  $u > t$  gives,

$$Y_u = y_t(u) + \int_t^u \kappa(u-s) \sqrt{V_s} dW_s \quad (4)$$

where

$$y_t(u) := \mathbb{E}_t [Y_u] = \bar{Y} + \int_{-\infty}^t \kappa(u-s) \sqrt{V_s} dW_s.$$

Then,  $y_t(u)$  is a martingale and

$$dy_t(u) = \kappa(u-t) \sqrt{V_t} dW_t.$$

## The forward variance curve

From the model definition, for  $u > t$ ,  $V_u = y_u(u)^2 + c$  so applying Itô's formula, we have

$$\begin{aligned} \xi_t(u) &= \mathbb{E}_t [y_u(u)^2] + c \\ &= y_t(u)^2 + \mathbb{E}_t \left[ \int_t^u d\langle y_t(u) \rangle_s \right] + c \\ &= y_t(u)^2 + \int_t^u \xi_t(s) \kappa(u-s)^2 ds + c. \end{aligned} \quad (5)$$

Alternatively,

$$y_t(u)^2 = \xi_t(u) - \int_t^u \xi_t(s) \kappa(u-s)^2 ds - c,$$

- Thus  $y_t(u)$  may be easily imputed from the forward variance curve.

## The QR Heston forward variance curve

- From the model definition , for  $u > t$ ,  $V_u = (Z_u - b)^2 + c = y_u(u)^2 + c$  so applying Itô's Formula,

```
\begin{aligned}
\xi_t(u) &:= \mathbb{E}_t[V_u] = \mathbb{E}_t[y_u(u)^2] + c \\
&= y_t(u)^2 + \mathbb{E}_t[\int_t^u d\langle y_t(u) \rangle_s] + c \\
&= y_t(u)^2 + \int_t^u \xi_t(s) \kappa(u-s)^2 ds + c.
\end{aligned}
```

- Alternatively,

$$y_t(u)^2 = \xi_t(u) - \int_t^u \xi_t(s) \kappa(u-s)^2 ds - c,$$

so, in principle,  $y_t(u)$  may be easily imputed from the forward variance curve.

## $\xi_t(u)$ from $y_t(u)$

- We thus have a Wiener-Hopf equation for  $\xi_t(u)$  whose solution may be written as

$$\xi_t(u) = y_t(u)^2 + c + \int_t^u K(u-s) [y_t(s)^2 + c] ds.$$

- The Laplace transform (denoted as  $\mathcal{L}$ ) of the resolvent kernel  $K$  is given by

$$\mathcal{L}[K] = \frac{\mathcal{L}[\kappa^2]}{1 - \mathcal{L}[\kappa^2]}.$$

## Dynamics of forward variance

- Using that  $\xi_t(u)$  is a martingale, we may write the QR Heston model in forward variance form:

$$\begin{aligned} d\xi_t(u) &= 2 y_t(u) dy_t(u) + 2 \int_t^u y_t(s) dy_t(s) K(u-s) ds \\ &= 2 y_t(u) \kappa(u-t) \sqrt{V_t} dW_t + 2 \int_t^u y_t(s) \kappa(s-t) K(u-s) ds \sqrt{V_t} dW_t \\ &= -2 \left\{ \kappa(u-t) y_t(u) + \int_t^u y_t(s) \kappa(s-t) K(u-s) ds \right\} \frac{dS_t}{S_t}. \end{aligned}$$

- We see that  $dS_t$  and  $d\xi_t(u)$  are usually (but not always) anti-correlated in the QR Heston model.

## Features of the quadratic rough Heston model

- Like the rough Heston model, the QR Heston model has a microstructural foundation.
  - The ZHawkes model of [Blanc et al.]<sup>[2]</sup> which is in turn related to QUARCH (quadratic ARCH).
- Dynamics are approximately lognormal, consistent with observation.
- VIX smiles are positively sloped, again consistent with observation.
- Like the other models, very parsimonious.
  - $S_t$  and  $\xi_t(\cdot)$  are state variables.
  - Parameters are  $c$  and the parameters of the chosen kernel, such as  $H$  and  $\eta$ .

## Summary

- We have presented the three most popular rough volatility models.
    - In all of these models,  $S_t$  and  $\xi_t(\cdot)$  are state variables.
  - We showed how to estimate the forward variance curve.
- We formulated all three models in forward variance form.
- The models are all very parsimonious.
    - All of them fit the SPX surface remarkably well.
  - The rough Bergomi model is easy to simulate, with reasonable dynamics, but generates flat VIX smiles.
    - We showed how to estimate rough Bergomi parameters, and checked the resulting smiles.
  - Only the QR Heston model generates positively-sloped VIX smiles.

## References

1. ^ Christian Bayer, Peter Friz and Jim Gatheral, Pricing under rough volatility, *Quantitative Finance* **16**(6) 887-904 (2016).
2. ^ Pierre Blanc, Jonathan Donier, and Jean-Philippe Bouchaud, Quadratic Hawkes processes for financial prices, *Quantitative Finance* **17**(2), 171-188 (2017).
3. ^ Mikkel Bennedsen, Asger Lunde, and Mikko S. Pakkanen, Hybrid scheme for Brownian semistationary processes, *Finance and Stochastics* **21**(4) 931-965, (2017).
4. ^ Mikkel Bennedsen, Asger Lunde, and Mikko S. Pakkanen, Decoupling the short-and long-term behavior of stochastic volatility, *Journal of Financial Econometrics*, **20**(5) 961-1006, (2021).
  
5. ^ Omar El Euch, Masaaki Fukasawa, and Mathieu Rosenbaum, The microstructural foundations of leverage effect and rough volatility, *Finance and Stochastics* **22** 3241-280 (2018).
6. ^ Jim Gatheral, *The Volatility Surface: A Practitioner's Guide*, John Wiley and Sons, Hoboken, NJ (2006).
7. ^ Blanka Horvath, Aitor Muguruza and Mehdi Tomas, Deep learning volatility: a deep neural network perspective on pricing and calibration in (rough) volatility models, *Quantitative Finance* **21**(1) 11-27 (2021).
8. ^ Antoine Jacquier, Claude Martini, and Aitor Muguruza, On VIX Futures in the rough Bergomi model, *Quantitative Finance* **18**(1) 45-61 (2018).
  
9. ^ Thibault Jaisson and Mathieu Rosenbaum, Rough fractional diffusions as scaling limits of nearly unstable heavy tailed Hawkes processes, *The Annals of Applied Probability* **26**(5) 2860-2882 (2016).
  
10. ^ Ryan McCrickerd and Mikko S Pakkanen, Turbocharging Monte Carlo Pricing for the Rough Bergomi Model, *Quantitative Finance* **18**(11) 1877-1886 (2018).